

# Lecture Notes for **Machine Learning in Python**



Professor Eric Larson  
**Adaptive Neural Network Optimization**

# Class Logistics and Agenda

- Agenda:
  - More optimization techniques
    - Review
    - Adaptive Learning Strategies
  - Town Hall for MLP

# Class Overview, by topic

Table Data  
Visualization

Numpy, Pandas, Seaborn  
Overviews with some in-depth discussion

Dimension  
Reduction and  
Image Processing

Scikit-learn, Scikit Image,  
Intuition only, Some mathematics

Linear and  
Logistic  
Regression

Numpy, Recreate API for Scikit-learn  
Detailed mathematics for simple optimization  
intuition for advanced optimization

Neural Networks  
and Back Prop.

Numpy  
Detailed mathematics for NN operations

Wide and Deep  
Networks

Convolutional  
Networks

Recurrent  
Networks

Keras, Tensorflow  
Intuition, Detailed implement.

Ethics in  
Language Models

ConceptNet  
Case studies

# Practical Initialization of Architectures

SQL programmers be like

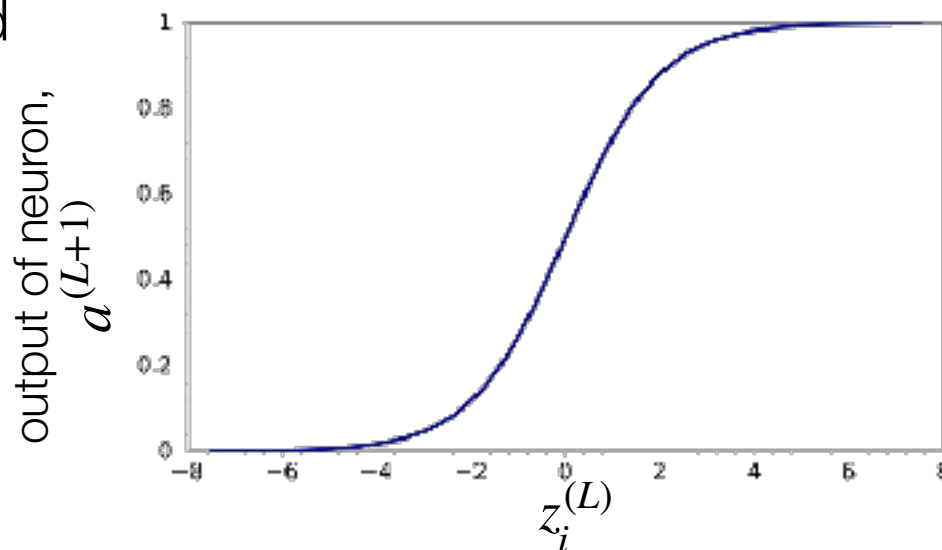


# Formative Self Test

- for adding Gaussian random variables, variances add together

$\mathbf{a}^{(L+1)} = \phi(\mathbf{W}^{(L)}\mathbf{a}^{(L)})$  assume each element of  $\mathbf{a}$  is Gaussian

- If you initialized the weights,  $\mathbf{W}$ , with too large variance, you would expect the output of the neuron,  $\mathbf{a}^{(L+1)}$ , to be:
  - A. saturated to “1”
  - B. saturated to “0”
  - C. could either be saturated to “0” or “1”
  - D. would not be saturated



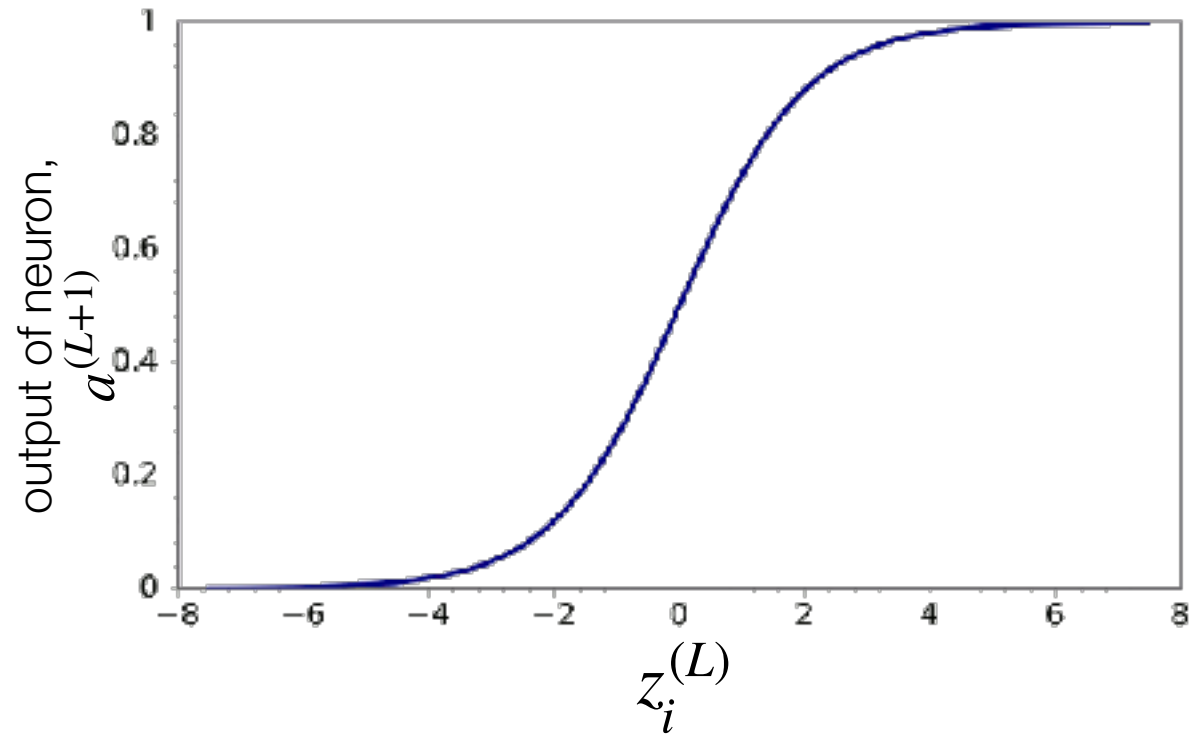
# Formative Self Test

- for adding Gaussian distributions, variances add together

$\mathbf{a}^{(L+1)} = \phi(\mathbf{W}^{(L)}\mathbf{a}^{(L)})$  assume each element of  $\mathbf{a}$  is Gaussian

- What is the derivative of a saturated sigmoid neuron?

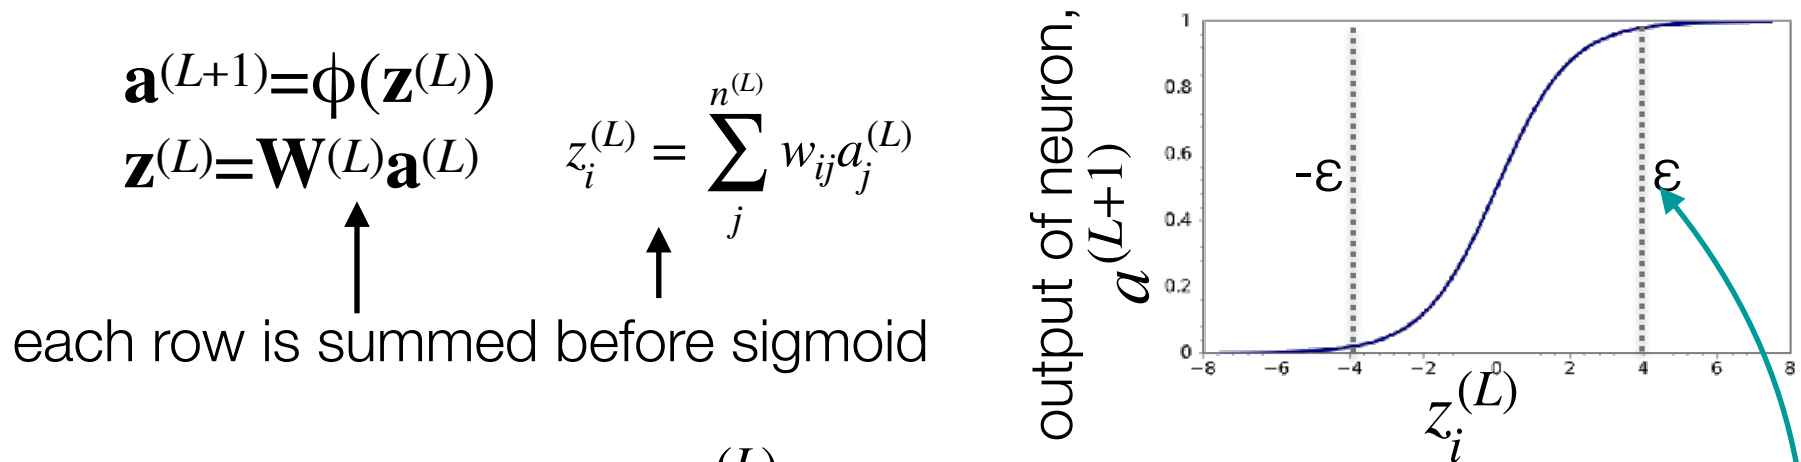
- A. zero
- B. one
- C.  $a \times (1 - a)$
- D. it depends



# Practical Implementation of Architectures

- **Weight initialization**

- try not to **saturate** your neurons right away!



want each  $z_i^{(L)}$  to be between  $-\epsilon < z_i^{(L)} < \epsilon$  for no saturation

**solution:** squash initial weights magnitude

- one choice: each element of  $\mathbf{W}$  selected from a Gaussian with **zero mean** and **specific standard deviation**

$$w_{ij}^{(L)} \approx \mathcal{N}\left(0, 4 \cdot \sqrt{\frac{1}{n^{(L)}}}\right)$$

For a sigmoid if,  $-\epsilon < z_i^{(L)} < \epsilon$  where  $\epsilon = 4$   
then  $a^{(L+1)}$  is well distributed  $[0, 1]$

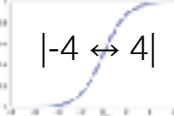
# Glorot Weight Initialization

Understanding the difficulty of training deep feedforward neural networks

Xavier Glorot JMLR 2010 Yoshua Bengio  
DIRO, Université de Montréal, Montréal, Québec, Canada

**Goal:** We should not saturate **feedforward** or **back propagated** variance

Relate variance of current layer to variance in  $z$ , so  $\sigma(z_i^{(L)})$  isn't saturated

  $|-4 \leftrightarrow 4|$  *try not to saturate*  $z_i^{(L)} = \sum_j^{n^{(L)}} w_{ij} a_j^{(L)}$  break down feed forward by multiply in  $i^{th}$  row

$$\text{Var}[z_i^{(L)}] = \sum_j^{n^{(L)}} \underbrace{E[w_{ij}]^2 \text{Var}[a_j^{(L)}] + \text{Var}[w_{ij}] E[a_j^{(L)}]^2 + \text{Var}[w_{ij}] \text{Var}[a_j^{(L)}]}_{0, \text{ if uncorrelated}} \quad \text{assume i.i.d. expand variance calc}$$

$$\text{Var}[z_i^{(L)}] = \sum_j^{n^{(L)}} \text{Var}[w_{ij}] \text{Var}[a_j^{(L)}] = n^{(L)} \underbrace{\text{Var}[w_{ij}] \text{Var}[a_j^{(L)}]}_{\approx 1} = n^{(L)} \text{Var}[w_{ij}]$$

$\text{Std}[z_i^{(L)}] = \sqrt{n^{(L)}} \cdot \text{Std}[w_{ij}]$   
*Want to keep  $\sim 4$*

$$\text{Std}[z_i^{(L)}] = 4 = \sqrt{n^{(L)}} \cdot \text{Std}[w_{ij}]$$

$$\text{Std}[w_{ij}] = 4 \cdot \sqrt{\frac{1}{n^{(L)}}}$$

$$w_{ij}^{(L)} \sim \mathcal{N}\left(0, 4 \cdot \sqrt{\frac{1}{n^{(L)}}}\right) \quad \begin{array}{l} \text{forward} \\ \text{from sigmoid} \end{array}$$



# Glorot Weight Initialization

$$\text{Std}[z_i^{(L)}] = 4 = \sqrt{n^{(L)}} \cdot \text{Std}[w_{ij}]$$

$$w_{ij}^{(L)} \sim \mathcal{N}\left(0, 4 \cdot \sqrt{\frac{1}{n^{(L)}}}\right) \text{ forward from sigmoid}$$

$$\mathbf{v}^{(L)} = \mathbf{a}^{(L)}(1 - \mathbf{a}^{(L)})\mathbf{W}^{(L)} \cdot \mathbf{v}^{(L+1)}$$

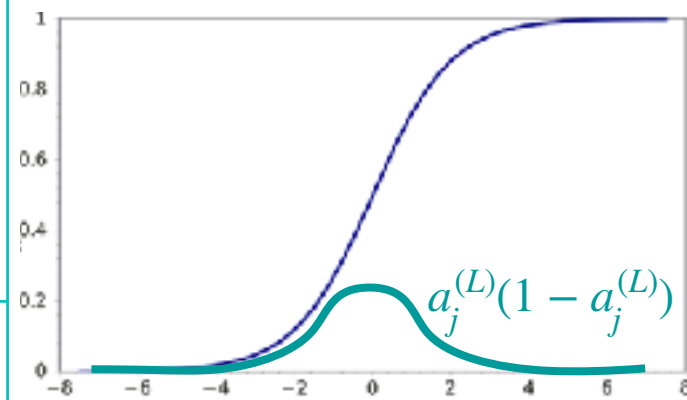
want to keep variance of  $\mathbf{v}$  stable  
magnitude  $\rightarrow$  stable gradient

Similar calculation for back prop.

$$\text{Var}[v_i^{(L)}] = n^{(L+1)}\text{Var}[w_{ij}]\text{Var}[v_j^{(L+1)} \cdot a_j^{(L)}(1 - a_j^{(L)})]$$

$$\text{Std}[v_i^{(L)}] = \sqrt{n^{(L+1)}} \cdot \text{Std}[w_{ij}] \cdot 0.25 \quad \text{want} = 1$$

$$w_{ij}^{(L)} \sim \mathcal{N}\left(0, 4 \cdot \sqrt{\frac{1}{n^{(L+1)}}}\right) \text{ backward from sensitivity}$$



$$w_{ij}^{(L)} \sim \mathcal{N}\left(0, 4 \cdot \sqrt{\frac{2}{n^{(L)} + n^{(L+1)}}}\right)$$

compromise

$$w_{ij}^{(L)} \sim U\left[\pm 4 \sqrt{\frac{6}{n^{(L)} + n^{(L+1)}}}\right]$$

if drawn from uniform dist.

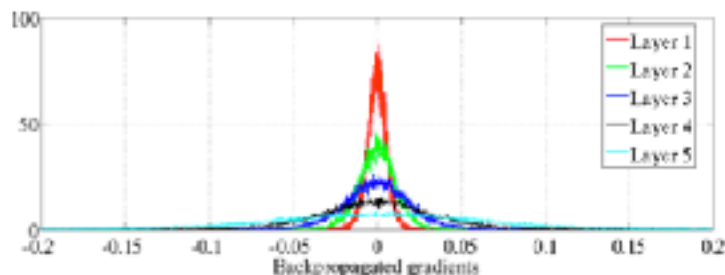
# Glorot Weight Initialization

## Understanding the difficulty of training deep feedforward neural networks

Xavier Glorot

DIRO, Université de Montréal, Montréal, Québec, Canada

Yoshua Bengio



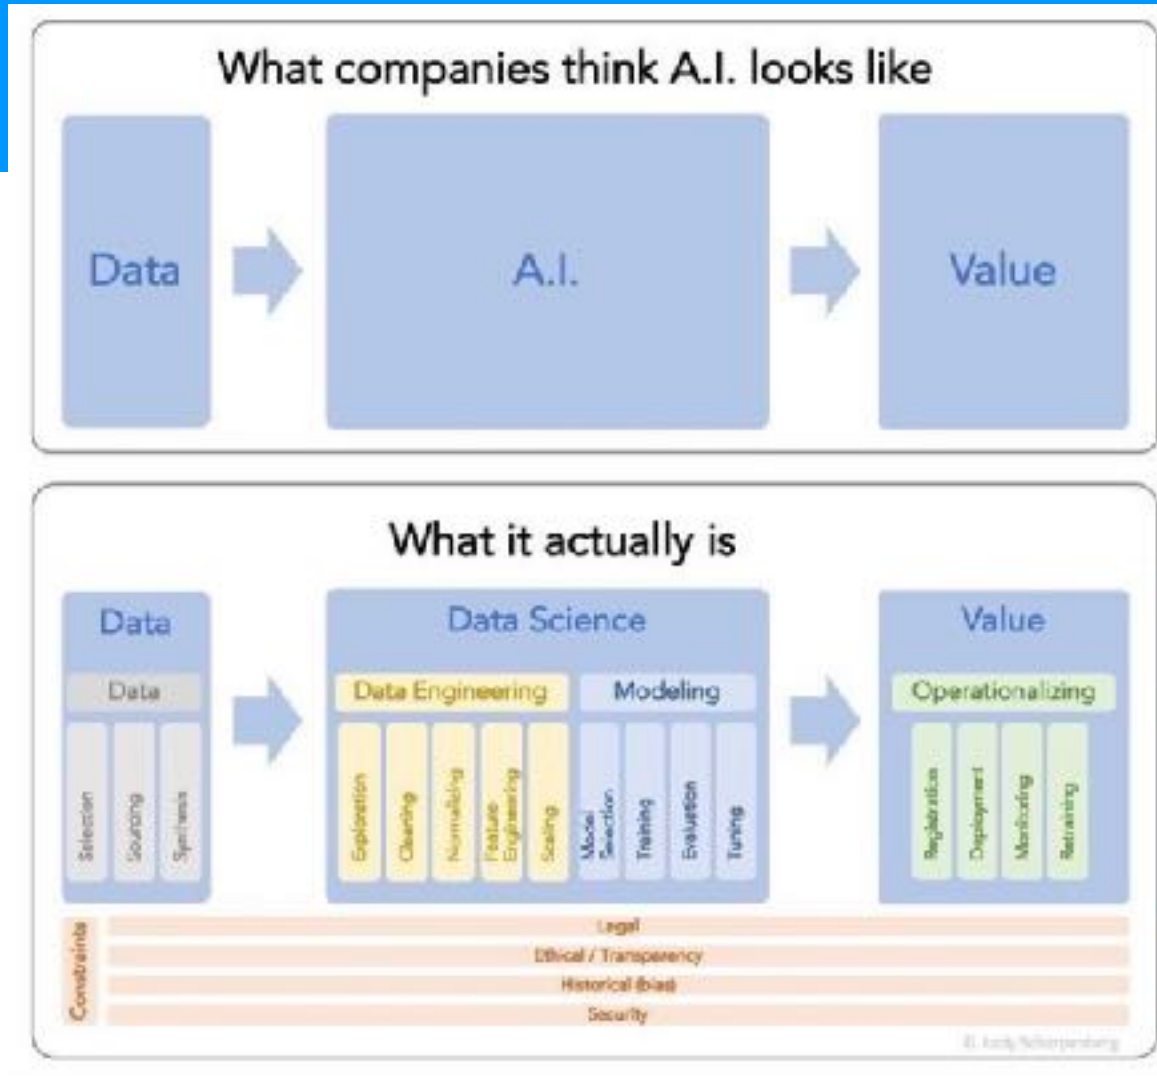
Starting gradient histograms  
per layer  
*standard initialization*

Figure 7: *Back-propagated gradients normalized histograms with hyperbolic tangent activation, with standard (top) vs normalized (bottom) initialization. Top: 0-peak decreases for higher layers.*

Smarter Weight Initialization

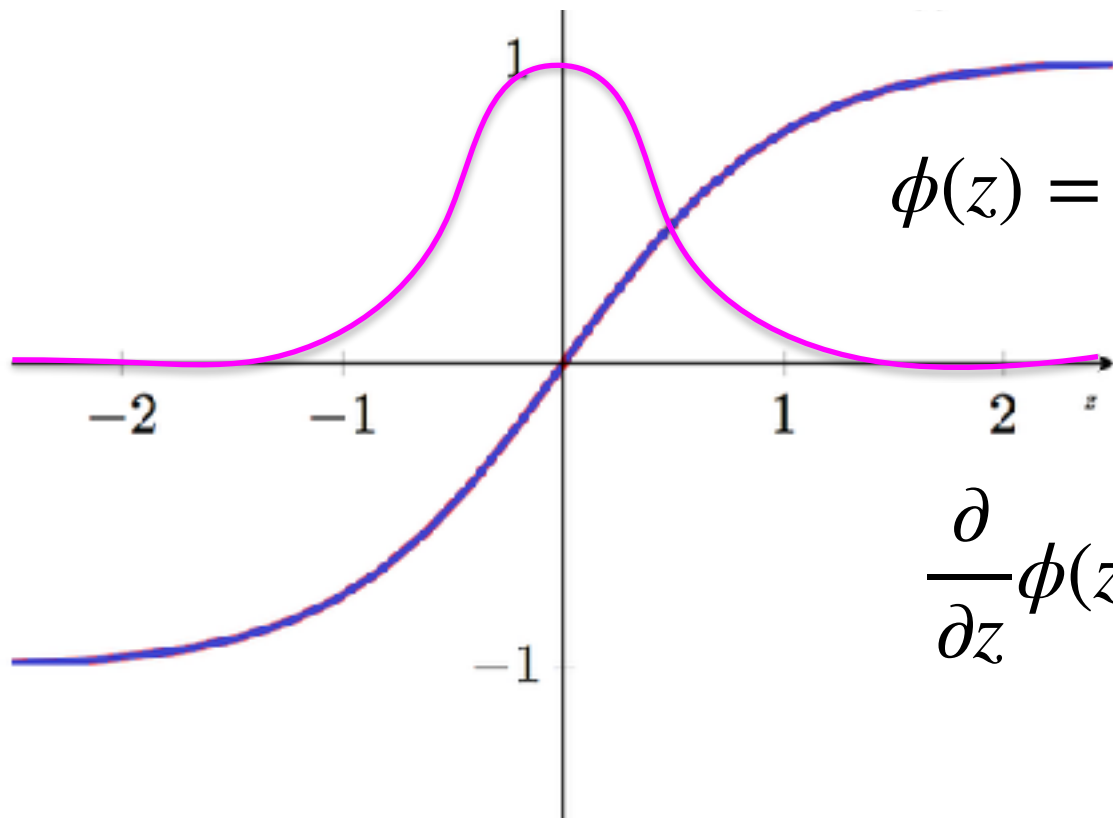


# Beyond Sigmoid: Other Activations



# New Activation: Hyperbolic Tangent

- Basically a sigmoid from -1 to 1

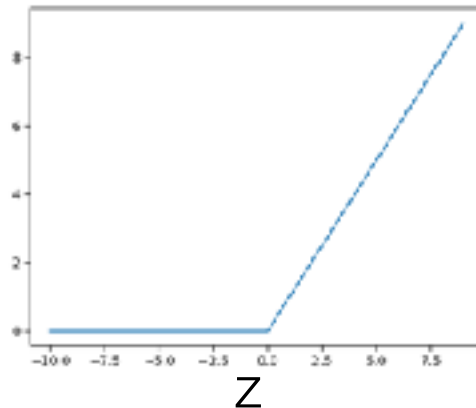
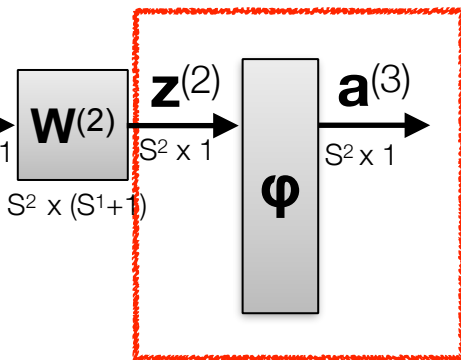


$$\phi(z) = \frac{\sinh(z)}{\cosh(z)} = \frac{e^z - e^{-z}}{e^z + e^{-z}}$$

$$\frac{\partial}{\partial z}\phi(z) = \text{sech}^2(z)$$

# New Activation: ReLU

- A new nonlinearity: **rectified linear units**

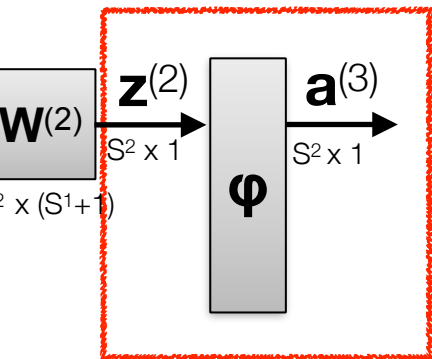


$$\phi(z) = \begin{cases} z, & \text{if } z > 0 \\ 0, & \text{else} \end{cases}$$

it has the advantage of **large gradients** and **extremely simple** derivative

$$\frac{\partial}{\partial z} \phi(z) = \begin{cases} 1, & \text{if } z > 0 \\ 0, & \text{else} \end{cases}$$

# Other Activation Functions



- Sigmoid Weighted Linear Unit **SiLU**
  - also called Swish
- Mixing of sigmoid,  $\sigma$ , and ReLU

$$\phi(z) = \sigma(z) \cdot z$$

Ramachandran P, Zoph B, Le QV. Swish: a Self-Gated Activation Function. arXiv preprint arXiv:1710.05941. 2017 Oct 16

Elfwing, Stefan, Eiji Uchibe, and Kenji Doya. "Sigmoid-weighted linear units for neural network function approximation in reinforcement learning." Neural Networks (2018).

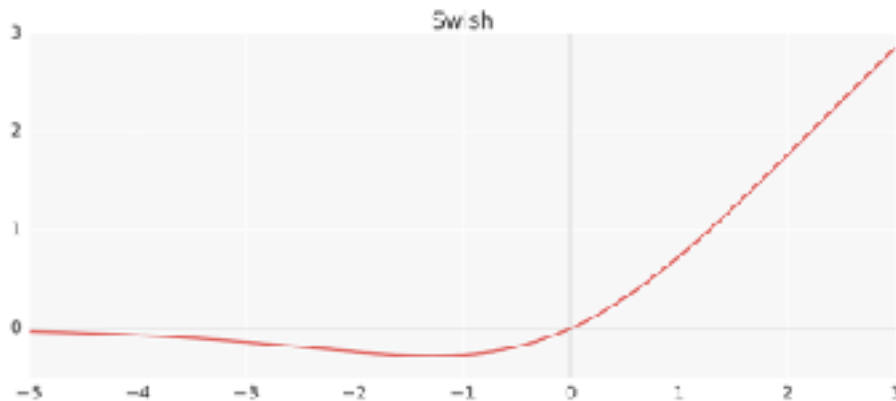


Figure 1: The Swish activation function.

$$\frac{\partial \phi(z)}{\partial z} = \frac{\partial}{\partial z} \sigma(z) \cdot z$$

$$= z \cdot \left[ \frac{\partial}{\partial z} \sigma(z) \right] + \sigma(z) \cdot \left[ \frac{\partial}{\partial z} z \right]$$

$$= z \cdot \sigma(z)(1 - \sigma(z)) + \sigma(z)$$

$$= z \cdot \sigma(z) + \sigma(z) \cdot (1 - z \cdot \sigma(z))$$

$$= \phi(z) + \sigma(z) \cdot (1 - \phi(z))$$

# Glorot and He Initialization

We have solved this assuming the activation output is in the range -4 to 4 (for a sigmoid) and assuming that we use Gaussian for sampling.

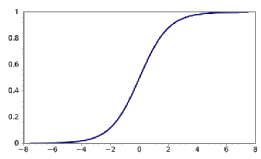
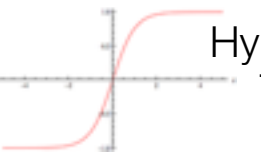
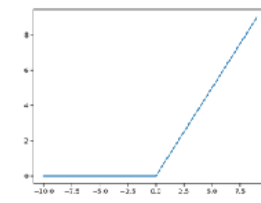
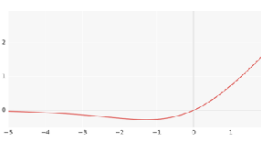
This range is different depending on the activation and assuming Gaussian or Uniform sampling.

	Uniform	Gaussian
Tanh	$w_{ij}^{(L)} \sim \sqrt{\frac{6}{n^{(L)} + n^{(L+1)}}}$	$w_{ij}^{(L)} \sim \sqrt{\frac{2}{n^{(L)} + n^{(L+1)}}}$
Sigmoid	$w_{ij}^{(L)} \sim 4\sqrt{\frac{6}{n^{(L)} + n^{(L+1)}}}$	$w_{ij}^{(L)} \sim 4\sqrt{\frac{2}{n^{(L)} + n^{(L+1)}}}$
ReLU SiLU	$w_{ij}^{(L)} \sim \sqrt{2}\sqrt{\frac{6}{n^{(L)} + n^{(L+1)}}}$	$w_{ij}^{(L)} \sim \sqrt{2}\sqrt{\frac{2}{n^{(L)} + n^{(L+1)}}}$

Summarized by Glorot and He



# Activations Summary

	Definition	Derivative	Weight Init (Uniform Bounds)
 <p>Sigmoid</p>	$\phi(z) = \frac{1}{1 + e^{-z}}$	$\nabla \phi(z) = a(1 - a)$	$w_{ij}^{(L)} \sim \pm 4 \sqrt{\frac{6}{n^{(L)} + n^{(L+1)}}}$
 <p>Hyperbolic Tangent</p>	$\phi(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}}$	$\nabla \phi(z) = \frac{4}{(e^z + e^{-z})^2}$	$w_{ij}^{(L)} \sim \pm \sqrt{\frac{6}{n^{(L)} + n^{(L+1)}}}$
 <p>ReLU</p>	$\phi(z) = \begin{cases} z, & \text{if } z > 0 \\ 0, & \text{else} \end{cases}$	$\nabla \phi(z) = \begin{cases} 1, & \text{if } z > 0 \\ 0, & \text{else} \end{cases}$	$w_{ij}^{(L)} \sim \pm \sqrt{2} \sqrt{\frac{6}{n^{(L)} + n^{(L+1)}}}$
 <p>SiLU</p>	$\phi(z) = \frac{z}{1 + e^{-z}}$	$\nabla \phi(z) = \phi(z) + \sigma(z) \cdot (1 - \phi(z))$	

## 08. Practical\_NeuralNets.ipynb

ReLU Nonlinearities  
Important for deep networks



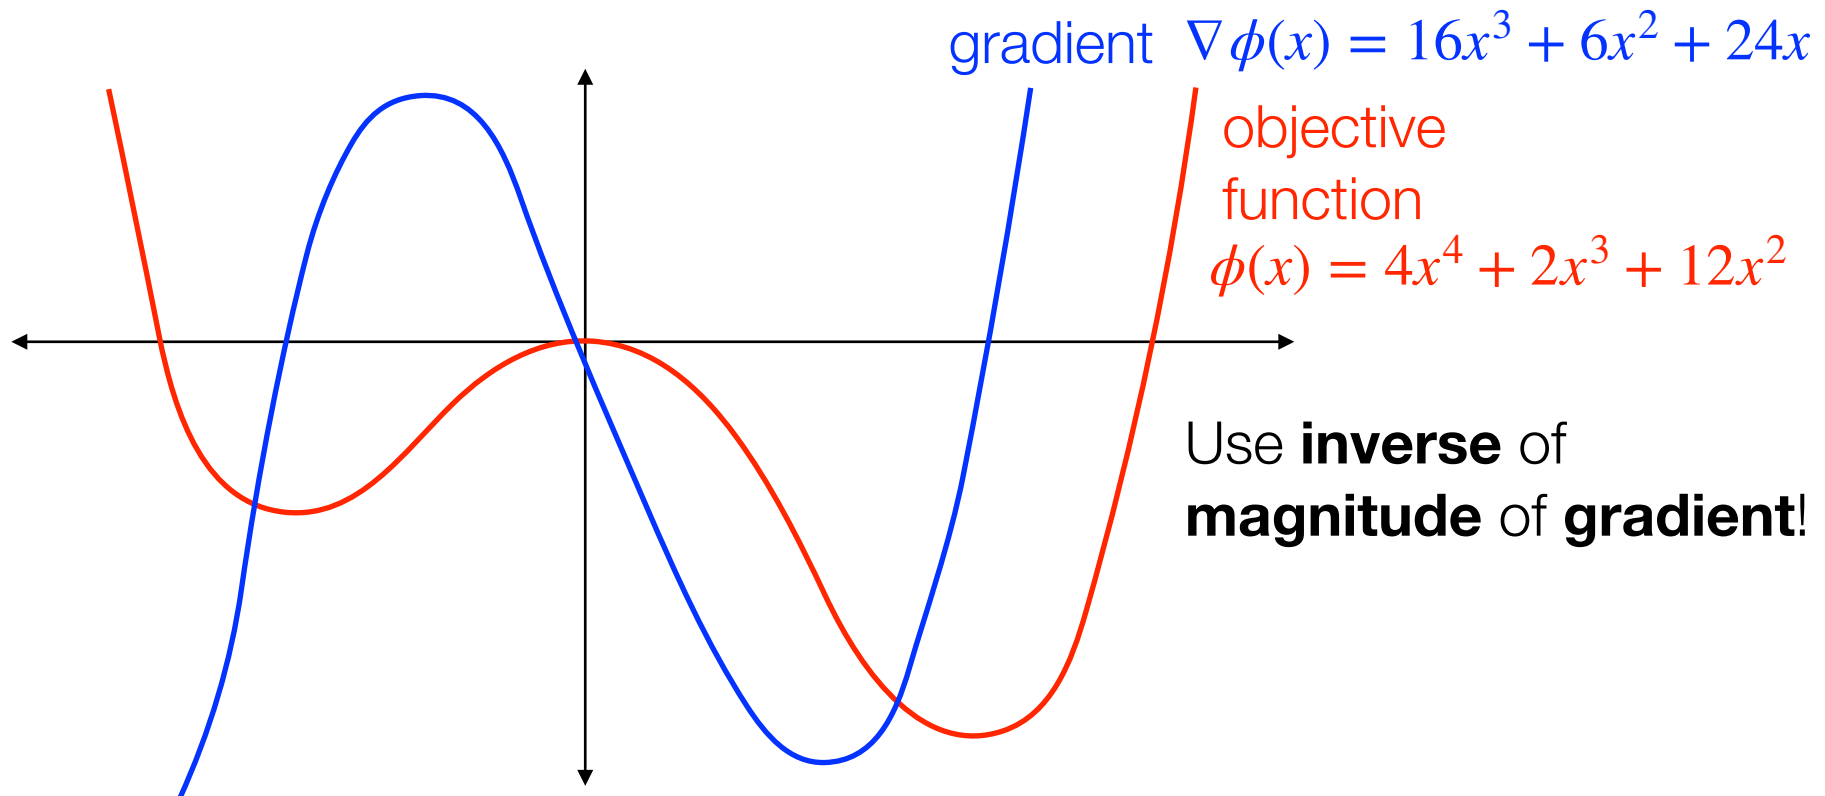
# More Adaptive Optimization



Going beyond  
changing the learning rate

# Be adaptive based on Gradient Magnitude?

- Decelerate down regions that are steep
- Accelerate on plateaus

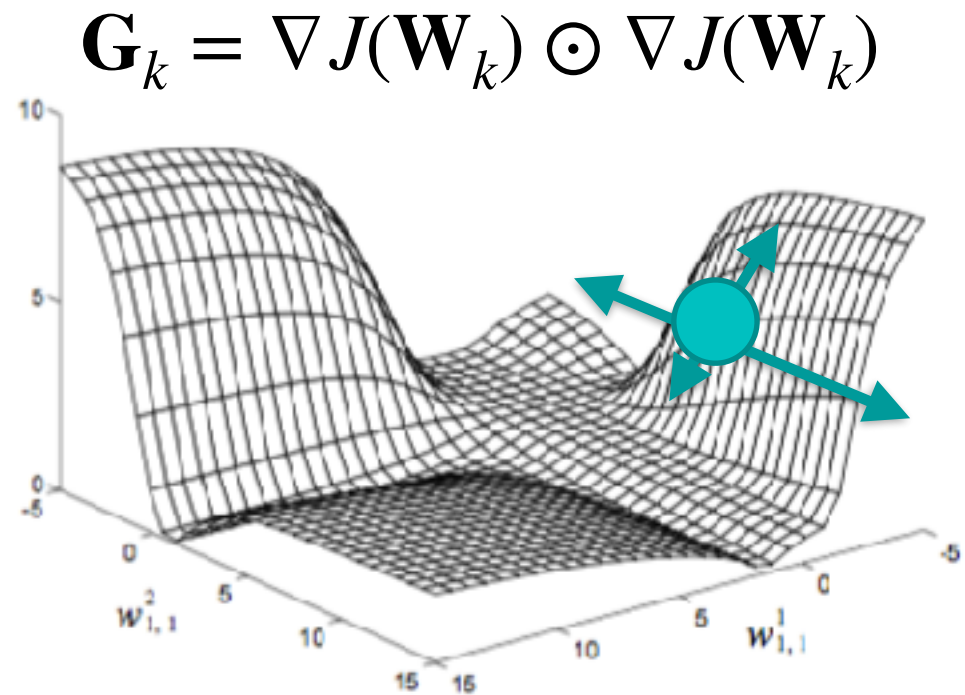
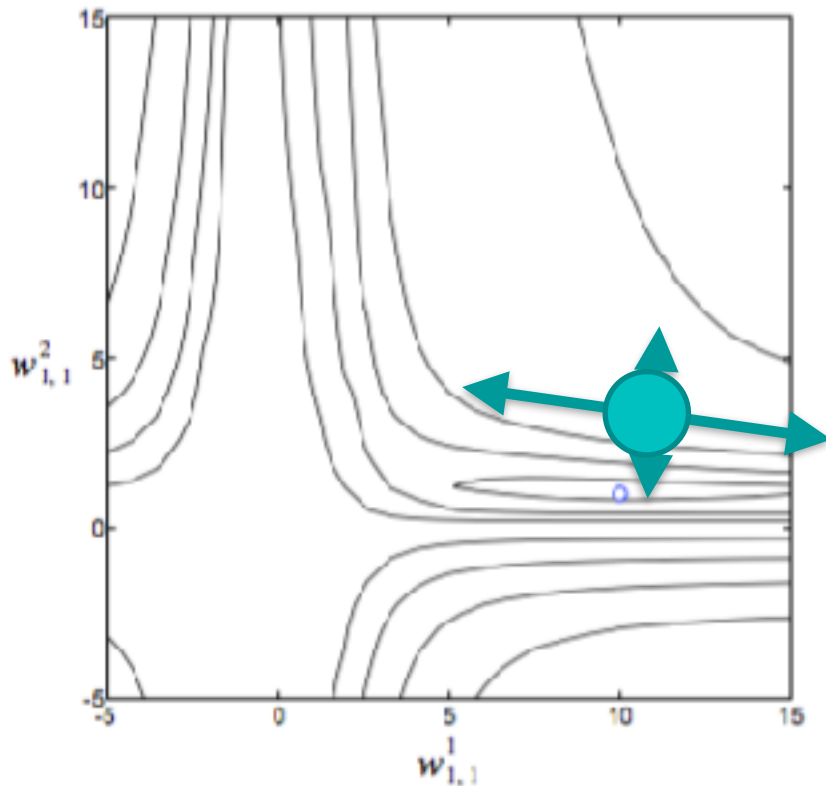


**Momentum:** be robust to **abrupt changes** in **steepness** (accumulate inverse magnitudes)

# Be adaptive based on Gradient Magnitude?

Inverse magnitude of gradient in multiple directions?

$$\mathbf{W}_{k+1} \leftarrow \mathbf{W}_k + \eta \frac{1}{\sqrt{\mathbf{G}_k + \epsilon}} \odot \nabla J(\mathbf{W}_k)$$



# Common Adaptive Strategies $\mathbf{W}_{k+1} = \mathbf{W}_k - \eta \cdot \rho_k$

Adjust each element of gradient by the steepness

- AdaGrad  
all operations are per element  
$$\rho_k = \frac{1}{\sqrt{\mathbf{G}_k + \epsilon}} \odot \nabla J(\mathbf{W}_k) \quad \text{where} \quad \mathbf{G}_k = \gamma \cdot \mathbf{G}_{k-1} + \nabla J(\mathbf{W}_k) \odot \nabla J(\mathbf{W}_k)$$
- RMSProp  
all operations are per element  
$$\rho_k = \frac{1}{\sqrt{\mathbf{V}_k + \epsilon}} \odot \nabla J(\mathbf{W}_k) \quad \begin{aligned} \mathbf{G}_k &= \nabla J(\mathbf{W}_k) \odot \nabla J(\mathbf{W}_k) \\ \mathbf{V}_k &= \gamma \cdot \mathbf{V}_{k-1} + (1 - \gamma) \cdot \mathbf{G}_k \end{aligned}$$
- AdaDelta  
all operations are per element  
$$\rho_k = \frac{\mathbf{M}_k}{\sqrt{\mathbf{V}_k + \epsilon}} \quad \mathbf{M}_{k+1} = \gamma \cdot \mathbf{M}_k + (1 - \gamma) \cdot \nabla J(\mathbf{W}_k)$$
- AdaM  
 $\mathbf{G}$  updates with decaying momentum of  $J$  and  $J^2$
- NAdaM  
same as Adam, but with nesterov's acceleration

**None** of these are “**one-size-fits-all**” because the space of neural network **optimization varies** by problem, AdaM is **popular** but **not a panacea**

# Adaptive Momentum

All operations are element wise:

$$\beta_1 = 0.9, \beta_2 = 0.999, \eta = 0.001, \epsilon = 10^{-8}$$

$$k = 0, \mathbf{M}_0 = \mathbf{0}, \mathbf{V}_0 = \mathbf{0}$$

Published as a conference paper at ICLR 2015

ADAM: A METHOD FOR STOCHASTIC OPTIMIZATION

Diederik P. Kingma\*  
University of Amsterdam, OpenAI

Jimmy Lei Ba\*  
University of Toronto

**For each epoch:**

update iteration  $k \leftarrow k + 1$

get gradient  $\nabla J(\mathbf{W}_k)$

accumulated gradient  $\mathbf{M}_k \leftarrow \beta_1 \cdot \mathbf{M}_{k-1} + (1 - \beta_1) \cdot \nabla J(\mathbf{W}_k)$

accumulated squared gradient  $\mathbf{V}_k \leftarrow \beta_2 \cdot \mathbf{V}_{k-1} + (1 - \beta_2) \cdot \nabla J(\mathbf{W}_k) \odot \nabla J(\mathbf{W}_k)$

boost moments magnitudes  
(notice  $k$  in exponent)  $\hat{\mathbf{M}}_k \leftarrow \frac{\mathbf{M}_k}{(1 - [\beta_1]^k)} \quad \hat{\mathbf{V}}_k \leftarrow \frac{\mathbf{V}_k}{(1 - [\beta_2]^k)}$

update gradient, normalized  
by second moment  
similar to AdaDelta  $\mathbf{W}_k \leftarrow \mathbf{W}_{k-1} - \eta \cdot \frac{\hat{\mathbf{M}}_k}{\sqrt{\hat{\mathbf{V}}_k + \epsilon}}$

gradient with momentum  
squared magnitude normalizer

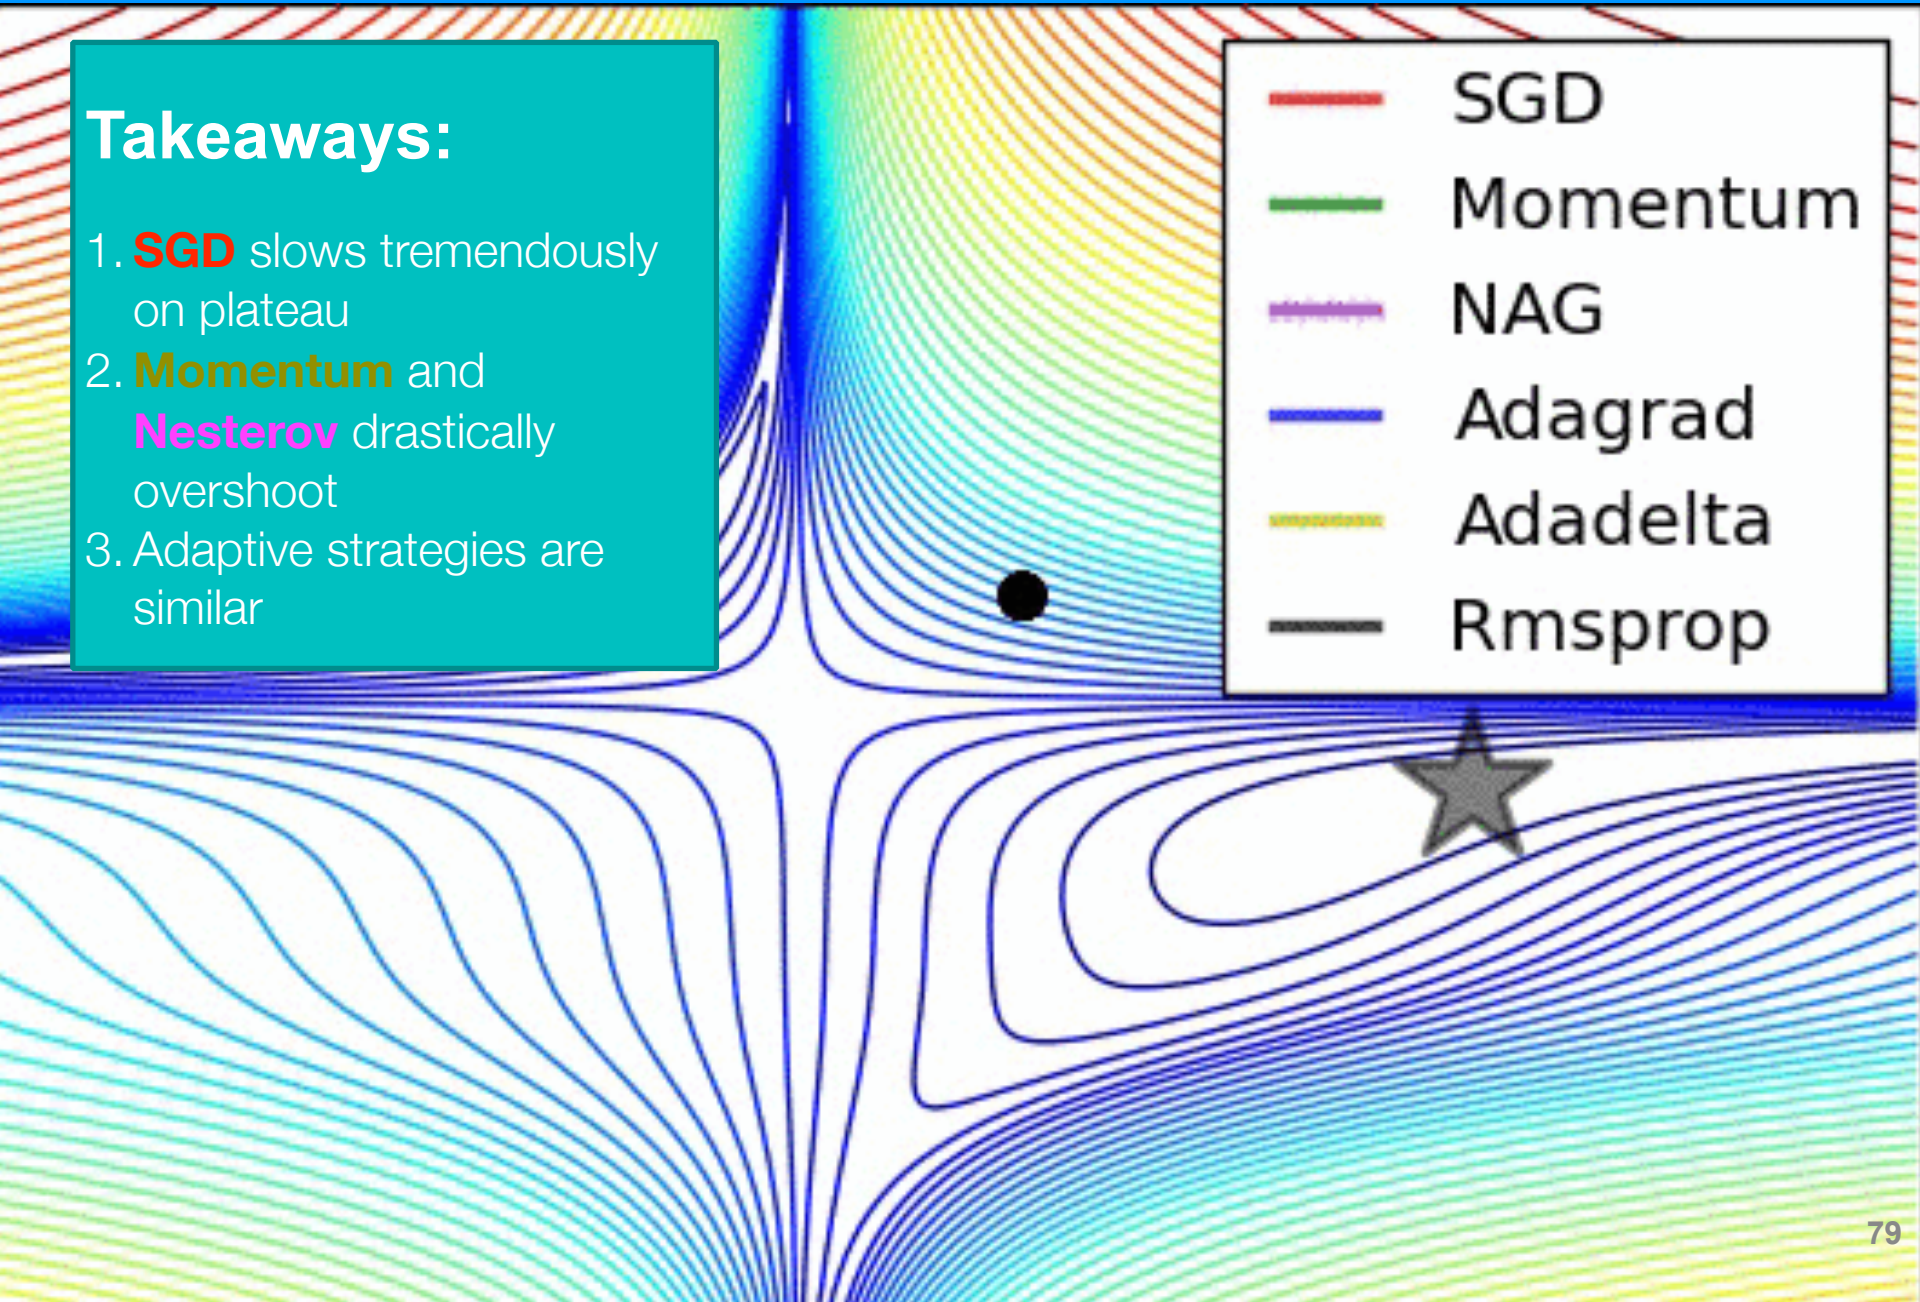
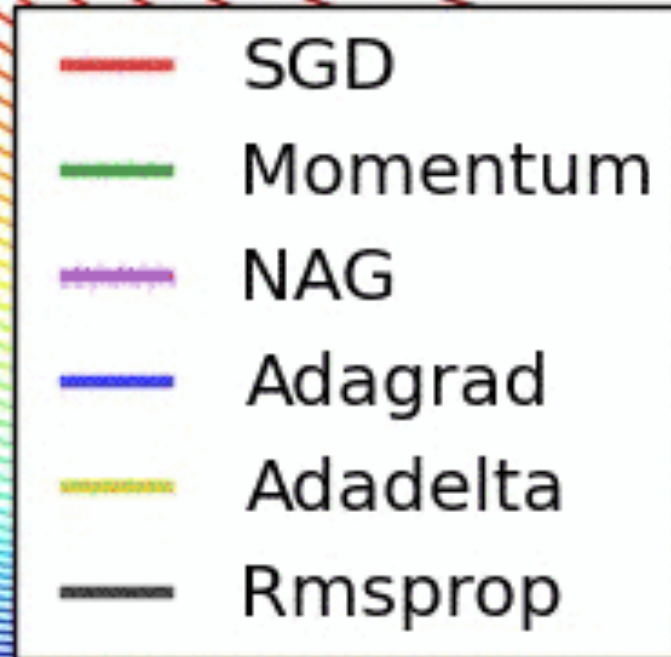


# Visualization of Optimization

<https://ruder.io/optimizing-gradient-descent/>

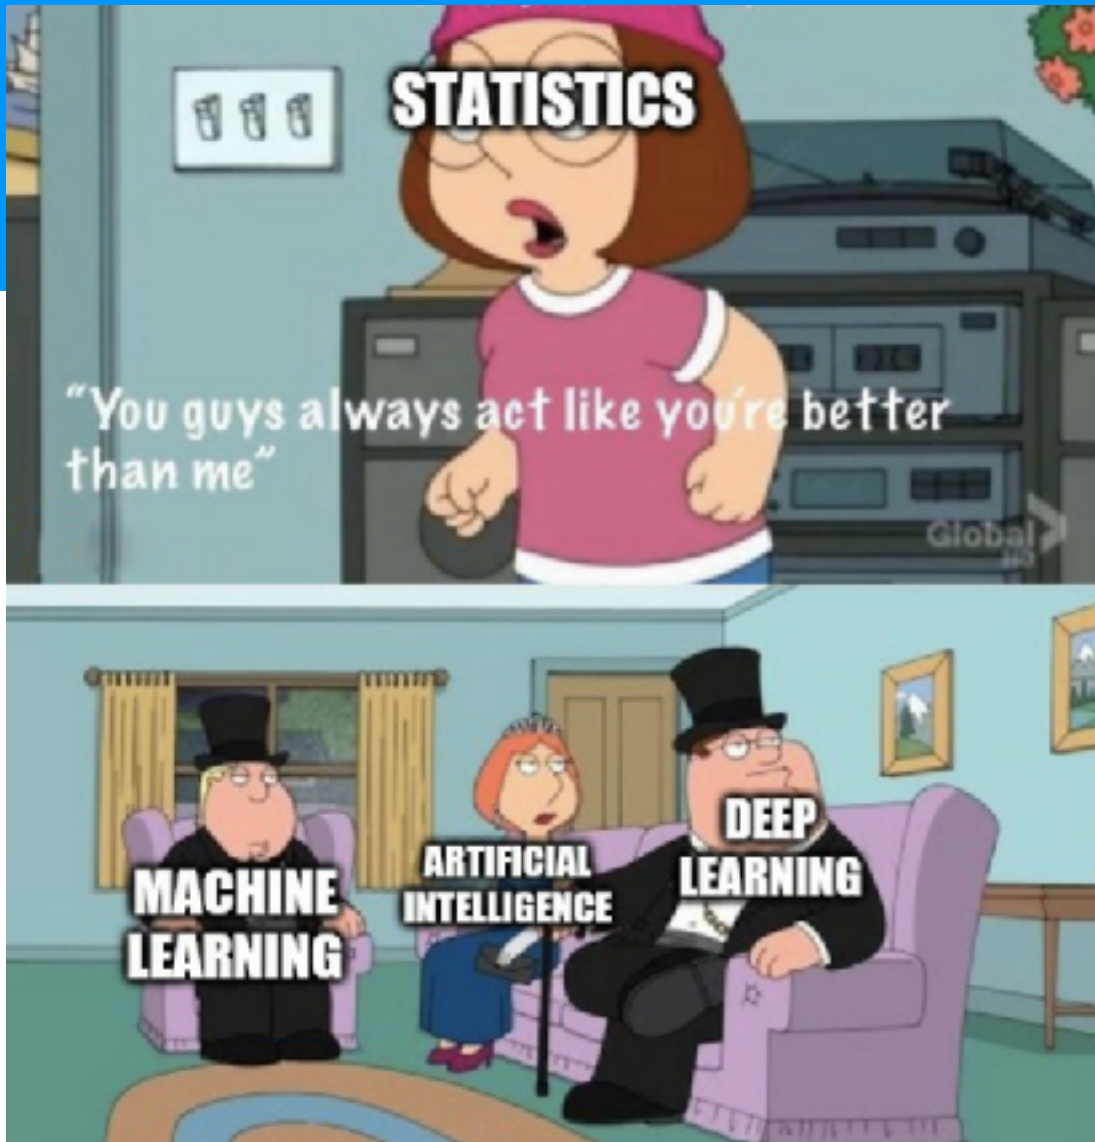
## Takeaways:

1. **SGD** slows tremendously on plateau
2. **Momentum** and **Nesterov** drastically overshoot
3. Adaptive strategies are similar





# Review



$$\mathbf{W}_{k+1} = \mathbf{W}_k - \rho_k$$

- Cross entropy  $\mathbf{A}^{(3)} - \mathbf{Y}$   
new final layer update

- Momentum  $\rho_k = \alpha \nabla J(\mathbf{W}_k) + \beta \nabla J(\mathbf{W}_{k-1})$

- Nesterov's Accelerated Gradient  $\rho_k = \underbrace{\beta \nabla J(\mathbf{W}_k + \alpha \nabla J(\mathbf{W}_{k-1}))}_{\text{step twice}} + \alpha \nabla J(\mathbf{W}_{k-1})$

- Mini-batching

← all data →

	batch 1	batch 2	batch 3	batch 4	batch 5	batch 6	batch 7	batch 8	batch 9
Epoch 1									
Epoch 2									
Epoch 3									
Epoch 4									
...									

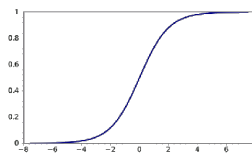
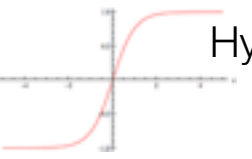
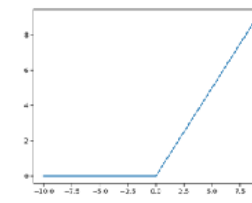
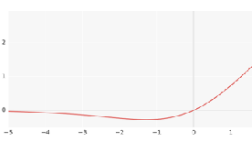
*shuffle ordering each epoch and update  $W$ 's after each batch*

- Learning rate adaptation (eta)

$$\eta_e = \eta_0^{(1+e \cdot \epsilon)}$$

$$\eta_e = \eta_0 \cdot d^{\lfloor \frac{e}{e_d} \rfloor}$$

# Review: Activations Summary

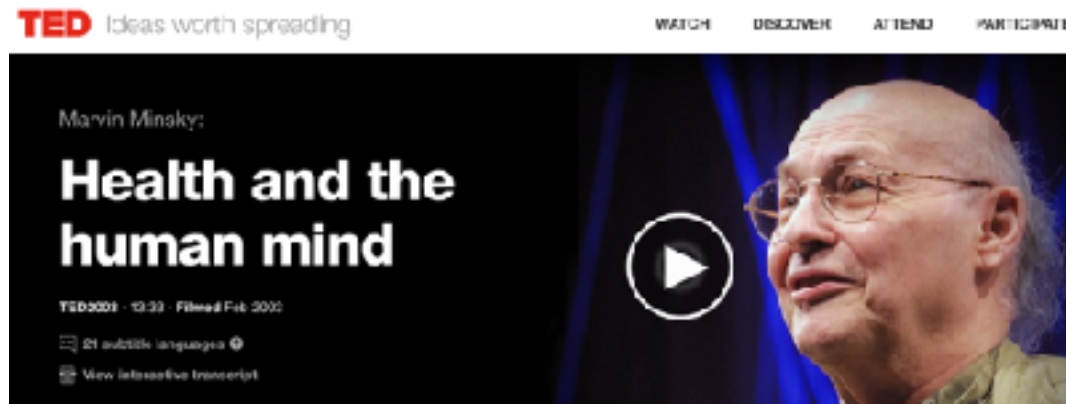
	Definition	Derivative	Weight Init (Uniform Bounds)
 <p>Sigmoid</p>	$\phi(z) = \frac{1}{1 + e^{-z}}$	$\nabla \phi(z) = a(1 - a)$	$w_{ij}^{(L)} \sim \pm 4 \sqrt{\frac{6}{n^{(L)} + n^{(L+1)}}}$
 <p>Hyperbolic Tangent</p>	$\phi(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}}$	$\nabla \phi(z) = \frac{4}{(e^z + e^{-z})^2}$	$w_{ij}^{(L)} \sim \pm \sqrt{\frac{6}{n^{(L)} + n^{(L+1)}}}$
 <p>ReLU</p>	$\phi(z) = \begin{cases} z, & \text{if } z > 0 \\ 0, & \text{else} \end{cases}$	$\nabla \phi(z) = \begin{cases} 1, & \text{if } z > 0 \\ 0, & \text{else} \end{cases}$	$w_{ij}^{(L)} \sim \pm \sqrt{2} \sqrt{\frac{6}{n^{(L)} + n^{(L+1)}}}$
 <p>SiLU</p>	$\phi(z) = \frac{z}{1 + e^{-z}}$	$\nabla \phi(z) = \phi(z) + \sigma(z) \cdot (1 - \phi(z))$	

# Revisiting Universality

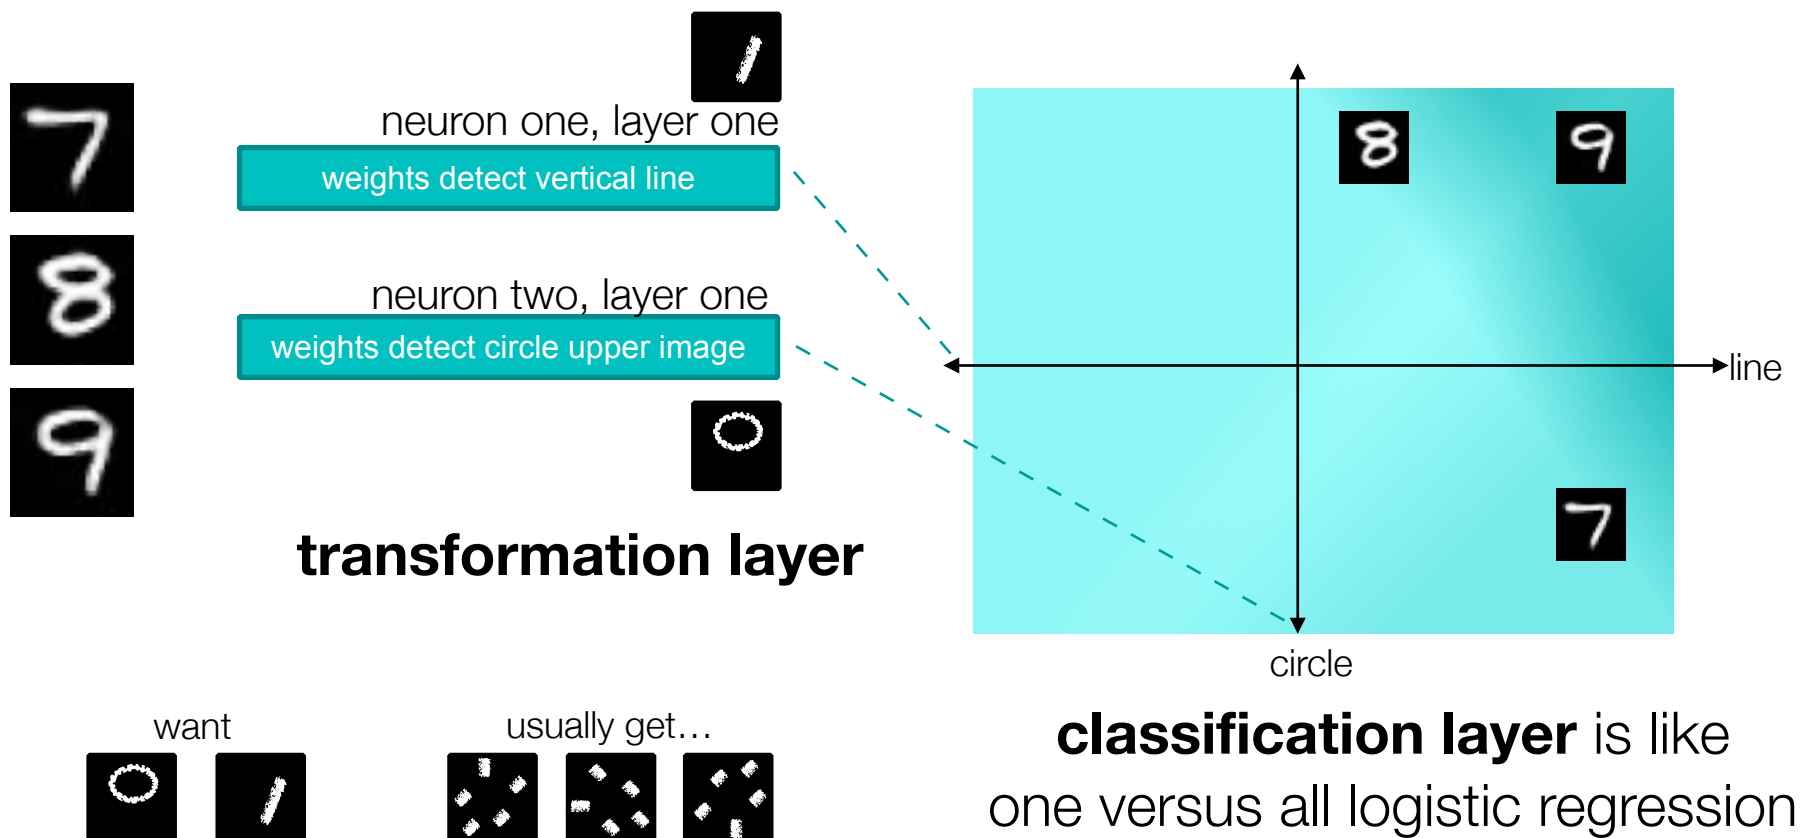
- Neural networks can separate any data through multiple layers. The true realization of Rosenblatt:

"Given an elementary  $\alpha$ -perceptron, a stimulus world  $W$ , and any classification  $C(W)$  for which a solution exists; let all stimuli in  $W$  occur in any sequence, provided that each stimulus must reoccur in finite time; then beginning from an arbitrary initial state, an error correction procedure will always yield a solution to  $C(W)$  in finite time..."

- **Universality:** No matter what function we want to compute, we know that there is a neural network which can do the job.



# Universality



- One nonlinear hidden layer with an output layer can perfectly train any problem with enough data, but might just be memorizing...
  - ... it might be better to have even more layers for decreased computation and generalizability

# End of Session

- Now: Lab 4 Town Hall
- Next Time: **Final Flipped Module!**
- Then: **Deep Learning**

# Town Hall

