

Lecture Notes for **Machine Learning in Python**

Professor Eric Larson
Lecture: RNN Demo

Lecture Agenda

- Logistics
 - RNNs due **During Finals Time**
- Recurrent Networks
 - *Overview*
 - *Problem Types*
 - *Embeddings*
 - *Types of RNNs*
 - **Demo A**
 - **CNNs and RNNs**
 - **Demo B**
 - **Ethics Case Study**
 - **Course Retrospective**

Class Overview, by topic

Table Data
Visualization

Numpy, Pandas, Seaborn
Overviews with some in-depth discussion

Dimension
Reduction and
Image Processing

Scikit-learn, Scikit Image,
Intuition only, Some mathematics

Linear and
Logistic
Regression

Numpy, Recreate API for Scikit-learn
Detailed mathematics for simple optimization
intuition for advanced optimization

Neural Networks
and Back Prop.

Numpy
Detailed mathematics for NN operations

Wide and Deep
Networks

Convolutional
Networks

Recurrent
Networks



Keras, Tensorflow
Intuition, Detailed implement.

Ethics in
Language Models

ConceptNet
Case studies


Last Time

- padding/clipping




Visualization

GloVe produces word vectors with a marked banded structure that is evident upon visualization:



Recurrent Networks

- basic RNN




$$h_t = \tanh(W_A (X_t @ h_{t-1}) + b_A)$$

$$P_t = \text{softmax}(W_P h_t + b_P)$$

<https://arxiv.org/pdf/1609.03126v1.pdf>

Self Test

- T/F: In Recurrent Neural Networks that are “rolled out”, each RNN cell can be run in parallel.
 - A. **True**, state vectors can be added later
 - B. **True**, but parallelization must use forward backward (like Viterbi)
 - C. **False**, state vectors must be found sequentially
 - D. **False**, input changes due to sequential nature of X_t

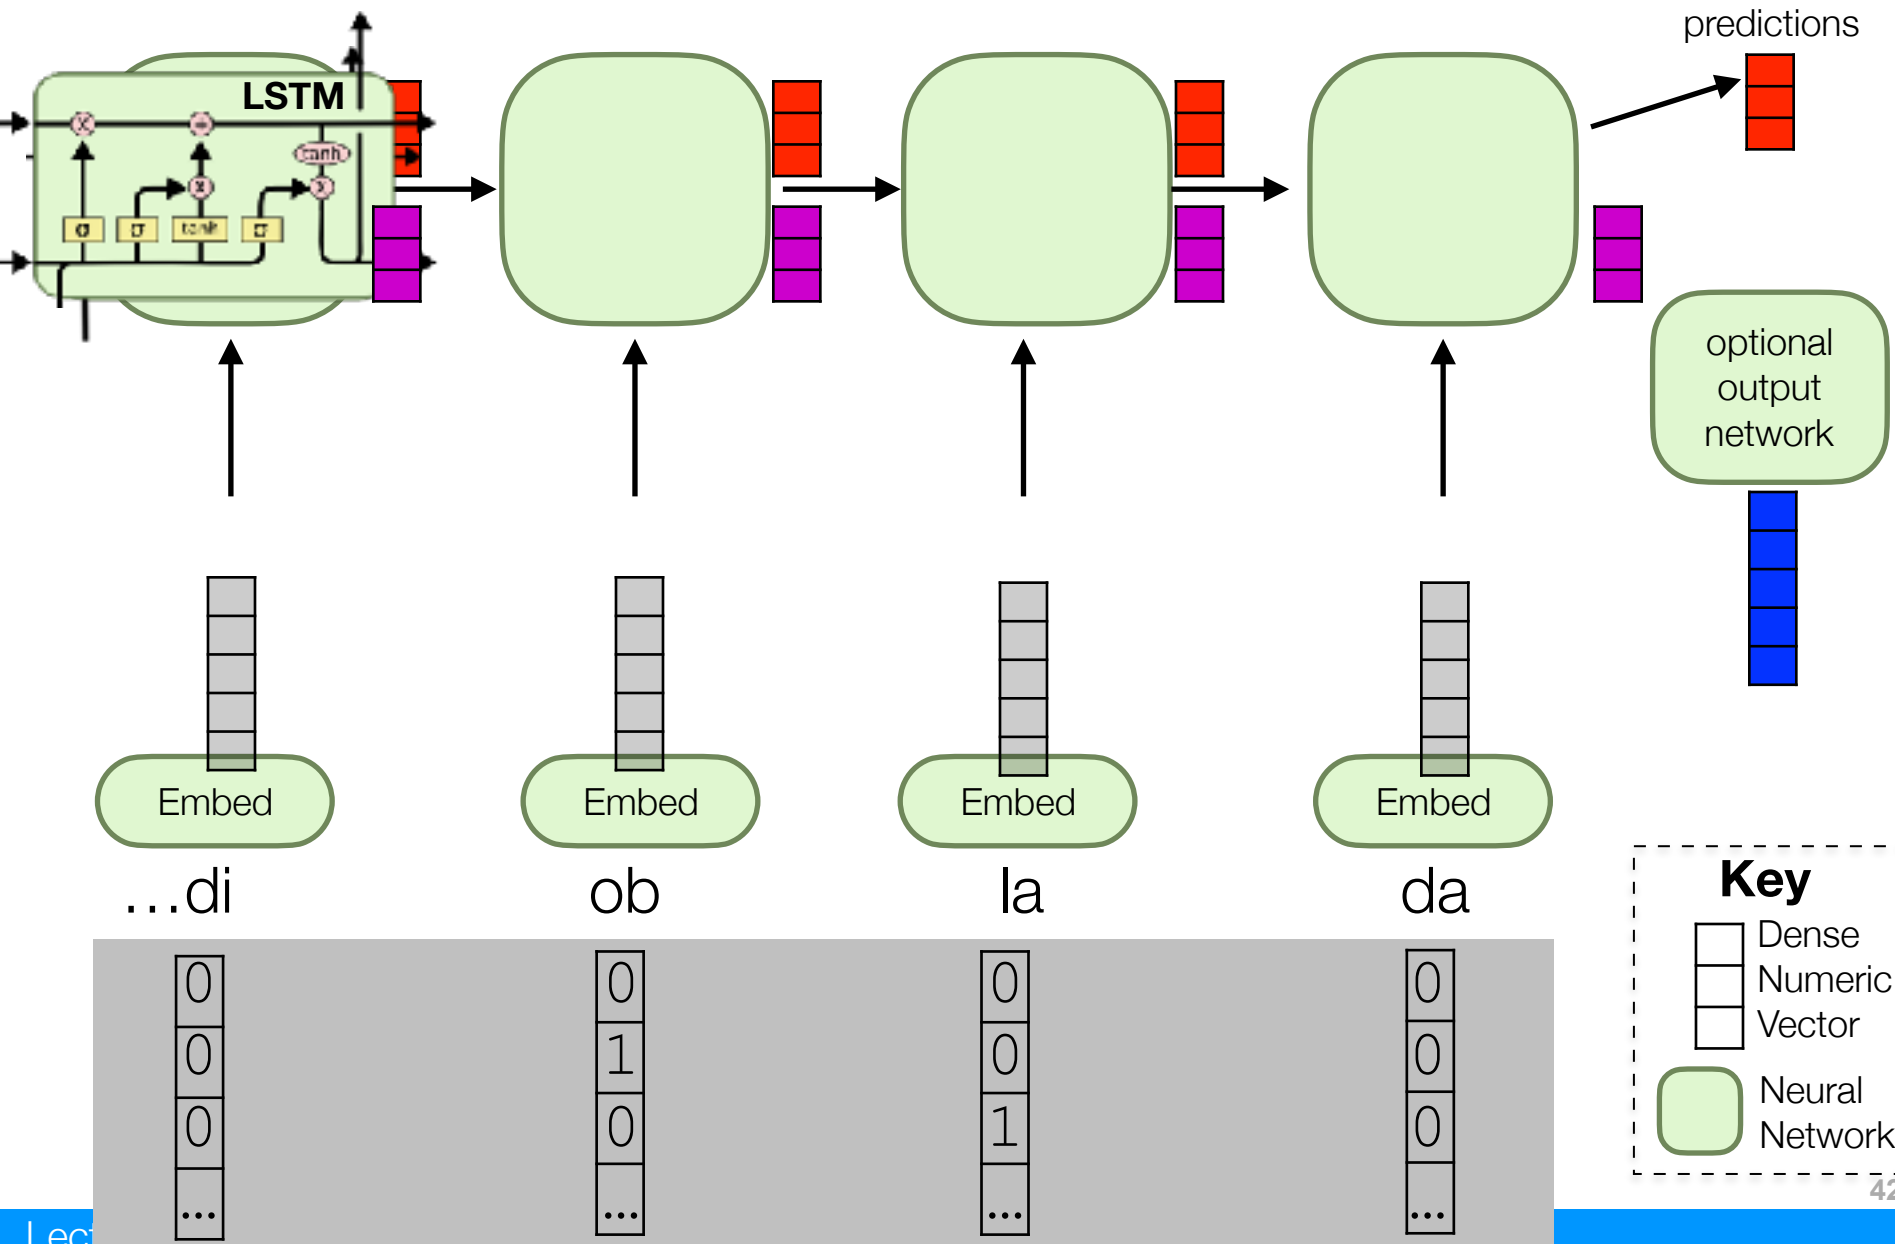


Commonly Used RNN Nodes



I like this version better.

General recurrent flow (many to one)



Recurrent Networks: GRUs

- gated recurrent units

Selectivity controls, gates (**0 or 1**)

$$r_t = \sigma(W_r s_{t-1} + U_r x_t + b_r)$$

$$z_t = \sigma(W_z s_{t-1} + U_z x_t + b_z)$$



past state

current input



selectively remember


with influence

$$\tilde{s}_t = \phi(W(r_t \odot s_{t-1}) + Ux_t + b)$$

$$s_t = z_t \odot s_{t-1} + (1 - z_t) \odot \tilde{s}_t$$

remember only past

OR remember with input



σ = sigmoid

\odot = elem. multiplication

Self Test


- What element of the GRU helps with vanishing and exploding gradients?
- A. derivative of σ
- B. no activation function
- C. derivative of ϕ
- D. ϕ

$$r_t = \sigma(W_r s_{t-1} + U_r x_t + b_r)$$


$$z_t = \sigma(W_z s_{t-1} + U_z x_t + b_z)$$

$$\tilde{s}_t = \phi(W(r_t \odot s_{t-1}) + Ux_t + b)$$

$$s_t = z_t \odot s_{t-1} + (1 - z_t) \odot \tilde{s}_t$$



Derivative of GRU



$$r_t = \sigma(W_r s_{t-1} + U_r x_t + b_r)$$

$$z_t = \sigma(W_z s_{t-1} + U_z x_t + b_z)$$

$$\tilde{s}_t = \phi(W(r_t \odot s_{t-1}) + U x_t + b)$$

$$s_t = z_t \odot s_{t-1} + (1 - z_t) \odot \tilde{s}_t$$

To back propagate, we need sensitivity of state vector, w.r.t previous state

Product Rule

Product Rule

$$\partial S_t / \partial S_{t-1} = (\partial z_t \times s_{t-1}) + (\partial s_{t-1} \times z_t) + \partial \tilde{s}_t - (\partial z_t \times \tilde{s}_t) - (\partial \tilde{s}_t \times z_t)$$

likely vanish

could vanish, depending on ϕ

likely vanish

could vanish, depending on ϕ

hard to vanish unless $z_t = 0$

Recurrent Networks: Gen 1 LSTM

- LSTM prototype

Selectivity controls (**gates, 0 or 1**)

$$o_t = \sigma(W_o s_{t-1} + U_o x_t + b_o)$$

$$i_t = \sigma(W_i s_{t-1} + U_i x_t + b_i)$$


$$f_t = \sigma(W_f s_{t-1} + U_f x_t + b_f)$$

$$\tilde{s}_t = \phi(W(o_t \odot s_{t-1}) + Ux_t + b)$$

selectively remember past with influence

$$s_t = f_t \odot s_{t-1} + i_t \odot \tilde{s}_t$$

selectively remember past with past weighted influence



Recurrent Networks: Gen 2 LSTM

- LSTM in TensorFlow

Selectivity controls (**gates, 0 or 1**)

$$i_t = \sigma(W_i h_{t-1} + U_i x_t + b_i)$$

$$o_t = \sigma(W_o h_{t-1} + U_o x_t + b_o)$$

$$f_t = \sigma(W_f h_{t-1} + U_f x_t + b_f)$$

explicit remembering state

$$\tilde{c}_t = \phi(W h_{t-1} + U x_t + b)$$

$$c_t = f_t \odot c_{t-1} + i_t \odot \tilde{c}_t$$

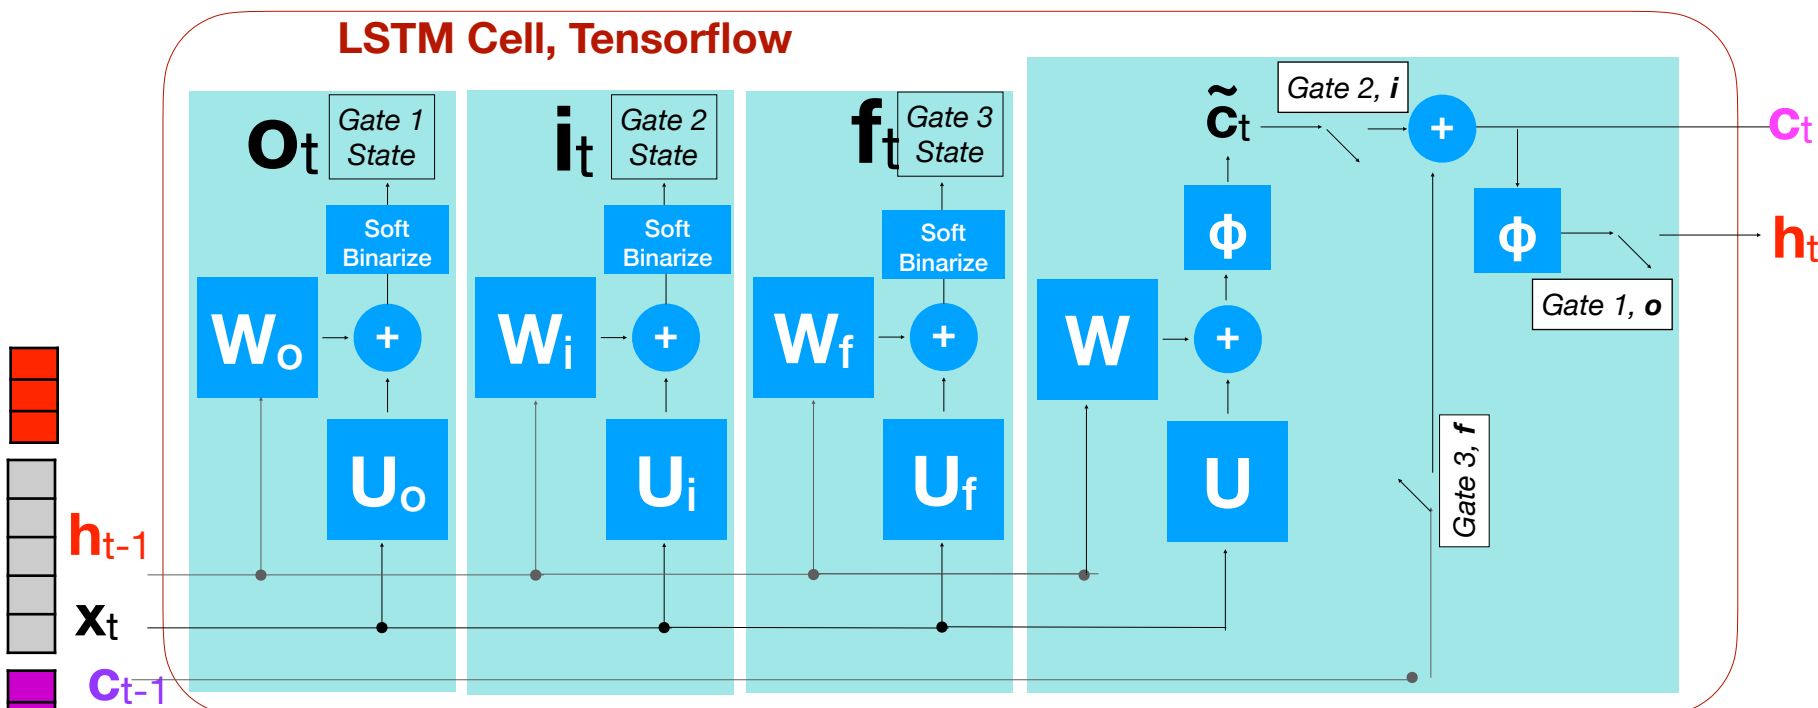
remember
previous state

update with
output, h_t

$$h_t = o_t \odot \phi(c_t)$$

get next h_t for
selecting gates

LSTM Cell, Tensorflow



LSTM Dropout


$$\begin{aligned}i_t &= \sigma(W_i h_{t-1} + U_i x_t + b_i) \\o_t &= \sigma(W_o h_{t-1} + U_o x_t + b_o) \\f_t &= \sigma(W_f h_{t-1} + U_f x_t + b_f)\end{aligned}$$

Recurrent
Dropout

Input
Dropout

The days of
training **without**
using **dropout** are
over.

LSTM Cell, Tensorflow



What to choose?

- There is no hard and fast rule
 - try both
 - basic LSTM has had great success
 - GRU also sometimes is easier to train
 - you will see many variations
 - peephole LSTM
 - hierarchical LSTM
 - and many more...

Many to one:
Simple RNNs
GRUs
LSTMs



More examples:

<https://github.com/tensorflow/tensorflow/tree/r0.11/tensorflow/examples/skflow>

<http://r2rt.com/recurrent-neural-networks-in-tensorflow-i.html>


<http://machinelearningmastery.com/sequence-classification-lstm-recurrent-neural-networks-python-keras/>

Seq2Seq:

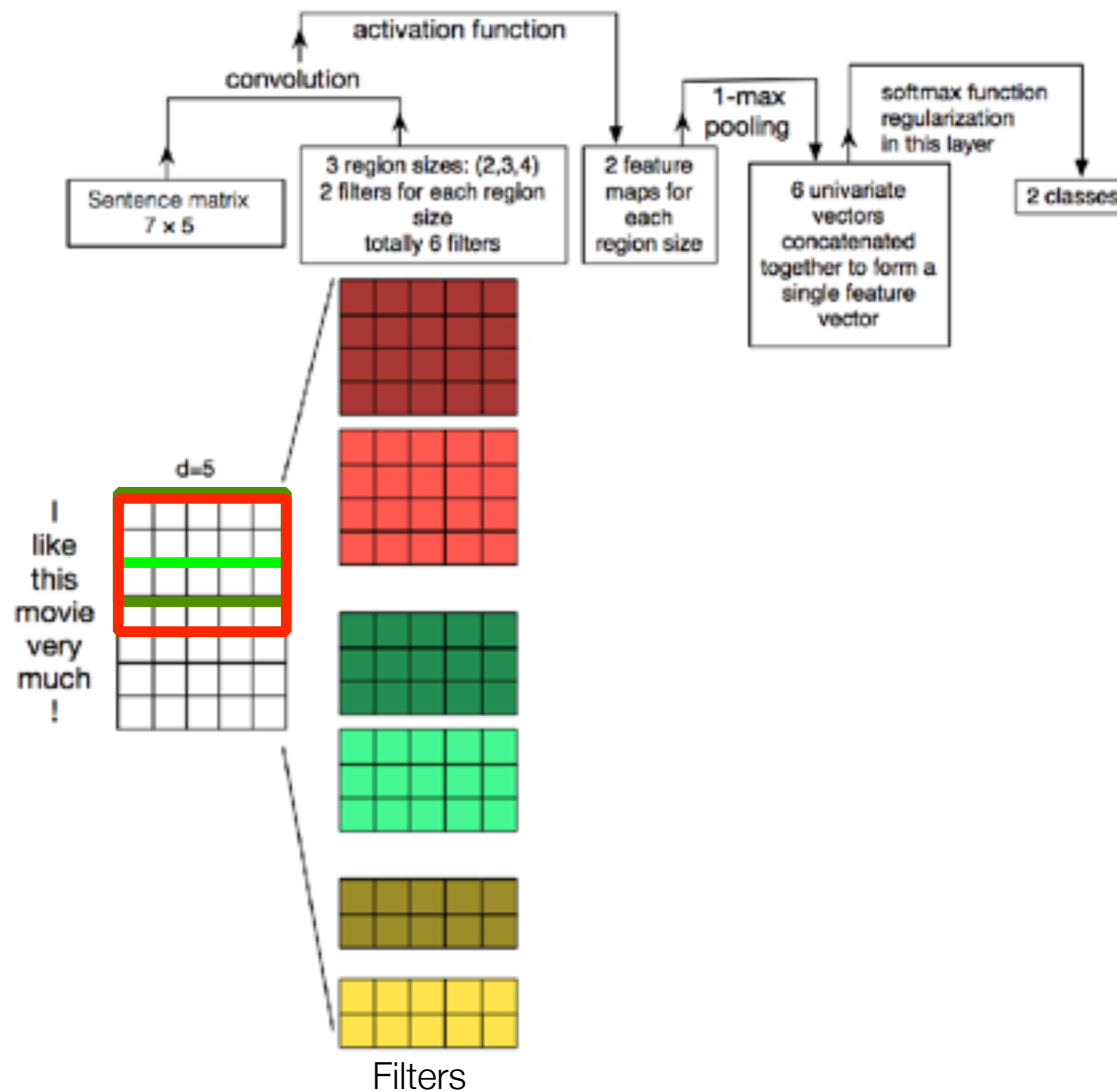
https://github.com/tensorflow/tensorflow/blob/r0.11/tensorflow/examples/skflow/neural_translation_word.py

CNNs for Sequences





CNNs with Multiple Region Sizes



Back to the CNN



More examples:

<http://www.wildml.com/2015/11/understanding-convolutional-neural-networks-for-nlp/>

Seq2Seq:

https://github.com/tensorflow/tensorflow/blob/r0.11/tensorflow/examples/skflow/neural_translation_word.py