

Lecture Notes for **Machine Learning in Python**



Professor Eric Larson
Keras: Wide and Deep Networks

Lecture Agenda

- Logistics:
 - CS 8321 in Spring
 - Grading and lab deadlines
- Review: Get out of the long winter...
- Introduction to TensorFlow
 - Tensors, Tf.Data
 - Deep APIs
- Wide and Deep Networks

Class Overview, by topic

Table Data
Visualization

Numpy, Pandas, Seaborn
Overviews with some in-depth discussion

Dimension
Reduction and
Image Processing

Scikit-learn, Scikit Image,
Intuition only, Some mathematics

Linear and
Logistic
Regression

Numpy, Recreate API for Scikit-learn
Detailed mathematics for simple optimization
intuition for advanced optimization

Neural Networks
and Back Prop.

Numpy
Detailed mathematics for NN operations

Wide and Deep
Networks

Convolutional
Networks

Recurrent
Networks

Keras, Tensorflow
Intuition, Detailed implement.

Ethics in
Language Models

ConceptNet
Case studies

Machine Learning Timeline


- **ImageNet competition occurs**
- **Second place:** 26.2% error rate
- **First place:**
 - From Hinton's lab, uses convolutional network with ReLU and dropout
 - 15.2% error rate
- Computer vision adopts deep learning with convolutional neural networks en masse



"I have had a hard time last fall so I really as she come pacin happ from skip a co Vision

Ever cleverer

Error rates on ImageNet Visual Recognition Challenge, %



Sources: ImageNet; Stanford Vision Lab


Economist.com



1949, Hebb's Law
Close neuron fire together



1960, Widrow & Hoff
Adaline Network




1986, Rumelhart & Hinton
Back-propagation



2003, Vapnik
Kernel SVMs



2012, Hinton, Fei-Fei Li
CNNs win ImageNet



Machine Learning Timeline (Neural Nets)

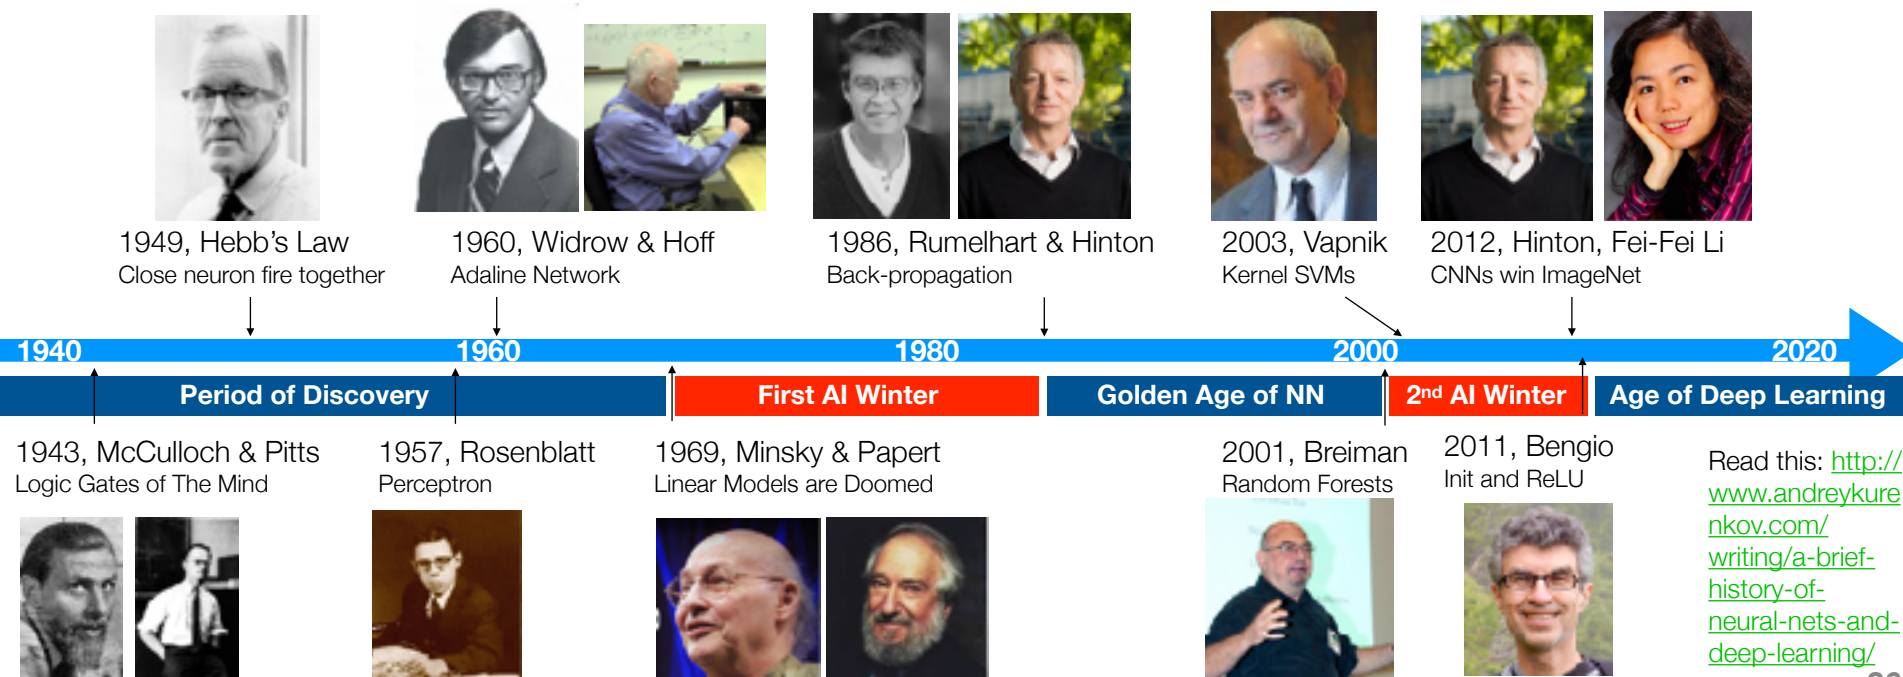
- 2012: Hinton Lab, Google, IBM, and Microsoft jointly publish paper, popularity for deep learning methods increases

Deep Neural Networks for Acoustic Modeling in Speech Recognition

[The shared views of four research groups]

[Geoffrey Hinton, Li Deng, Dong Yu, George E. Dahl, Abdel-rahman Mohamed, Navdeep Jaitly, Andrew Senior, Vincent Vanhoucke, Patrick Nguyen, Tara N. Sainath, and Brian Kingsbury]


<https://www.cs.toronto.edu/~gdahl/papers/deepSpeechReviewSPM2012.pdf>









Machine Learning Timeline (Neural Nets)

- 2013: Andrew Ng and Google (BrainTeam)
 - run unsupervised feature creation on YouTube videos (becomes computer vision benchmark)

The work resulted in unsupervised neural net learning of an unprecedented scale - 16,000 CPU cores powering the learning of a whopping 1 billion weights. The neural net was trained on Youtube videos, entirely without labels, and learned to recognize the most common objects in those videos.



A summary of the Deep Learning people:

					
Yoshua Bengio	Yann LeCun	Geoffrey Hinton	FeiFei Li	Andrew Ng	Daphne Koller
Stayed at Univ. Montreal Advises IBM	Heads Facebook AI Team	Univ. Toronto Google	Stanford (HAI) Former Chief Scien., AI/ML Google Cloud	Coursera Baidu Google	Stanford Founded Coursera MacArthur Genius



- Hinton: Restricted Boltzmann Machine, Deep autoencoder
- Bengio: neural language modeling.
- LeCun: Convolutional Neural Network
- NIPS, ICML, CVPR, ACL
- Google Brain, Deep Mind.
- FaceBook AI.

Made Deep Learning Instruction Accessible

doi:10.1088/nature14539

ing

Geoffrey Hinton⁴⁵

deep learning

Deep learning allows computational models that are composed of multiple processing layers to learn representations of data with multiple levels of abstraction. These methods have dramatically improved the state-of-the-art in speech recognition, visual object recognition, object detection and many other domains such as drug discovery and genomics. Deep learning discovers intricate structure in large data sets by using the backpropagation algorithm to indicate how a machine should change its internal parameters that are used to compute the representation in each layer from the representation in the previous layer. Deep convolutional nets have brought about breakthroughs in processing images, video, speech and

Credit for Deep Learning

Official ACM @TheOfficialACM

Yoshua Bengio, Geoffrey Hinton and Yann LeCun, the fathers of #DeepLearning, receive the 2018 #ACMTuringAward for conceptual and engineering breakthroughs that have made deep neural networks a critical component of computing today. bit.ly/2HVJtdV



Yoshua Bengio



Geoffrey Hinton



Yann LeCun

— citations_Foucault
— citations_Bourdieu
— citations_Altman



projected citations

Machine learning is the science of credit assignment. The machine learning community itself profits from proper credit assignment to its members. The inventor of an important method should get credit for inventing it. She may not always be the one who popularizes it. Then the popularizer should get credit for popularizing it (but not for inventing it). Relatively young research areas such

Review of Deep Learning History


- Up to this point: back propagation saved AI winter for NN (Hinton and others!)
- 80's, 90's, 2000's: convolutional networks for image processing start to get deeper
 - but back propagation no longer does great job at training them
- SVMs and Random Forests gain traction...
 - The second AI winter begins, research in NN plummets
- 2004: Hinton secures funding from CIFAR in 2004 Hinton rebrands: Deep Learning
- 2006: Auto-encoding and Restricted Boltzmann Machines
- 2007: Deep networks are more efficient when pre-trained
- 2009: GPUs decrease training time by 70 fold...
- 2010: Hinton's students go to internships with Microsoft, Google, and IBM, making their speech recognition systems faster, more accurate and deployed in only 3 months...
- 2012: Hinton Lab, Google, IBM, and Microsoft jointly publish paper, popularity sky-rockets for deep learning methods
- 2011-2013: Ng and Google run unsupervised feature creation on YouTube videos (becomes computer vision benchmark)
- 2012+: Pre-training is not actually needed, just solutions for vanishing gradients (like ReLU, SiLU, initializations, more data, GPUs)



1949, Hebb's Law
Close neuron fire together



1960, Widrow & Hoff
Adaline Network



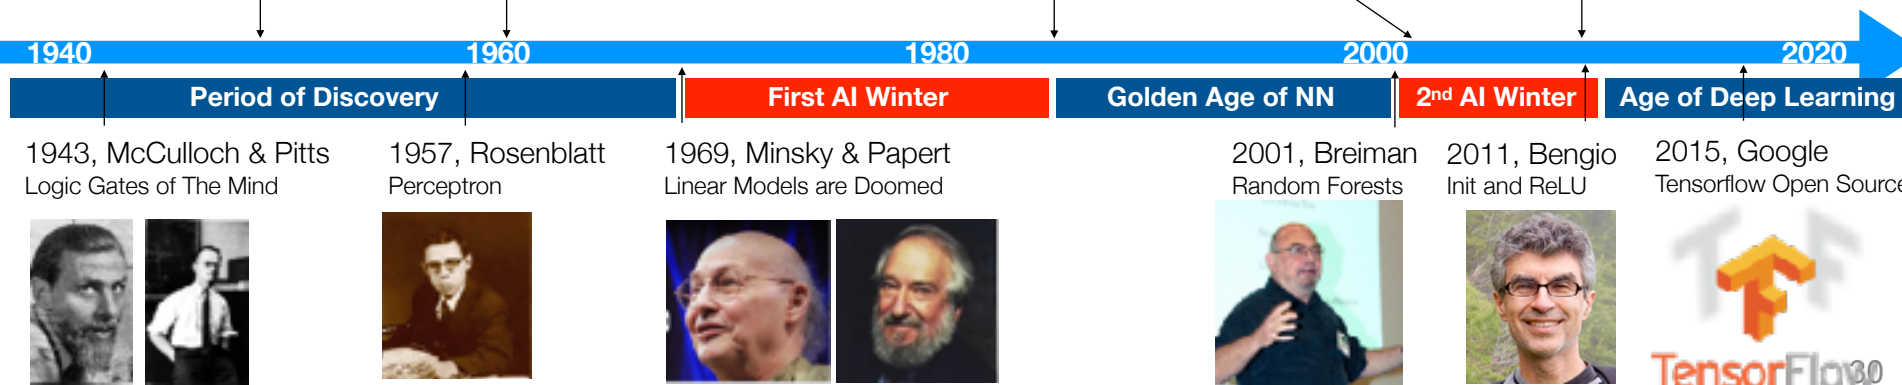
1986, Rumelhart & Hinton
Back-propagation



2003, Vapnik
Kernel SVMs



2012, Hinton, Fei-Fei Li
CNNs win ImageNet



Last Time


- Up to this point: back propagation saved AI winter for NN (Hinton and others!)
- 80's, 90's, 2000's: convolutional networks for image processing start to get deeper
 - but back propagation no longer does great job at training them
- SVMs and Random Forests gain traction...
 - The second AI winter begins, research in NN plummets
- 2004: Hinton secures funding from CIFAR in 2004 Hinton rebrands: Deep Learning
- 2006: Auto-encoding and Restricted Boltzmann Machines
- 2007: Deep networks are more efficient when pre-trained
- 2009: GPUs decrease training time by 70 fold...
- 2010: Hinton's students go to internships with Microsoft, Google, and IBM, making their speech recognition systems faster, more accurate and deployed in only 3 months...
- 2012: Hinton Lab, Google, IBM, and Microsoft jointly publish paper, popularity sky-rockets for deep learning methods
- 2011-2013: Ng and Google run unsupervised feature creation on YouTube videos (becomes computer vision benchmark)
- 2012+: Pre-training is not actually needed, just solutions for vanishing gradients (like ReLU, SiLU, initializations, more data, GPUs)



1949, Hebb's Law
Close neuron fire together



1960, Widrow & Hoff
Adaline Network



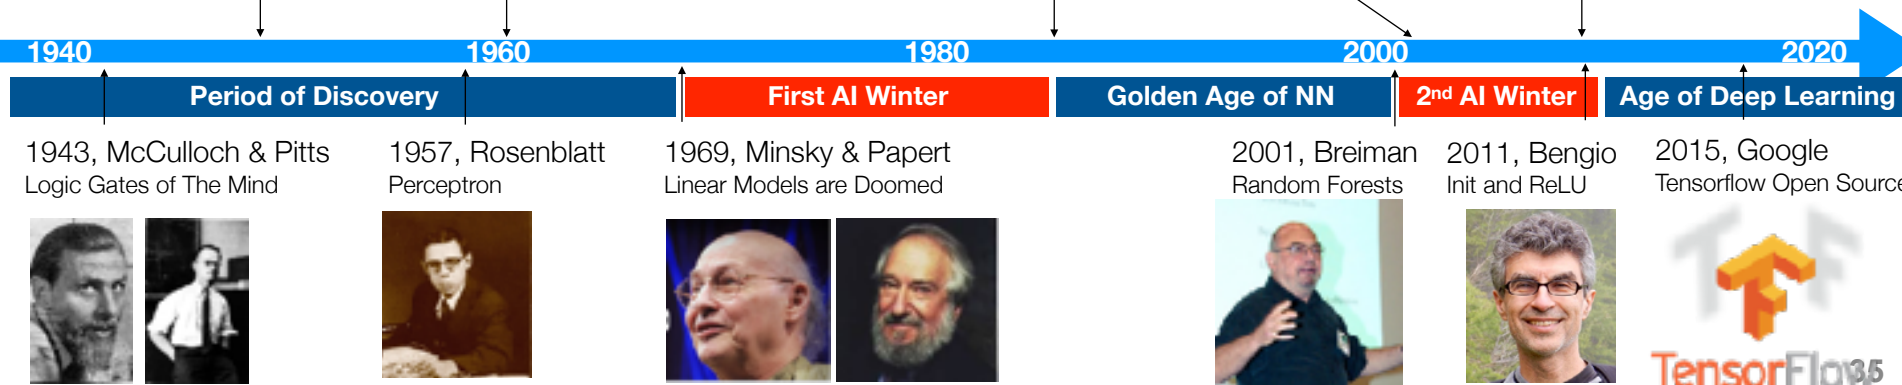
1986, Rumelhart & Hinton
Back-propagation



2003, Vapnik
Kernel SVMs



2012, Hinton, Fei-Fei Li
CNNs win ImageNet



TensorFlow

“Further discussion of it merely incumbers the literature and befogs the mind of fellow students.”

- 2007: **NIPS** program committee rejects a paper on deep learning by *al. et.* Hinton because they already accepted a paper on deep learning and two papers on the same topic would be excessive.
- ~2009: A reviewer tells Yoshua Bengio that papers about neural nets have no place in **ICML**.
- ~2010: A **CVPR** reviewer rejects Yann LeCun's paper even though it beats the state-of-the-art. The reviewer says that it tells us nothing about computer vision because everything is learned.




Options for Deep Learning Toolkits

1.  TensorFlow
2.  Keras
3.  PyTorch
4.  Caffe
5.  theano
6.  Apache MXNet
7.  Microsoft CNTK
8.  DL4J
9.  Caffe2
10.  Chainer
11.  fast.ai

Overview of Deep Learning frameworks adoption metrics over 2020

2022 Machine Learning & Data Science Survey by Kaggle: library usage (N=14,531)



Tensorflow

- Open sourced library from Google
- Second generation release from Google Brain
 - supported for Linux, Unix, Windows
 - Also works on Android/iOS
- Released November 9th, 2015
 - (this class first offered January 2016)



Programmatic creation

- Most toolkits use python to build a **computation graph** of operations
 - Build up computations
 - Execute computations
- **Most Toolkits Support:**
 - tensor creation
 - functions on tensors
 - automatic differentiation
- Tensors are just multidimensional arrays
 - like in Numpy
 - scalars (biases and constants)
 - vectors (e.g., input arrays)
 - 2D matrices (e.g., images)
 - 3D matrices (e.g., color images)
 - 4D matrices (e.g., batches of color images)

Tensor basic functions

- Easy to define operations on tensors

```
a = tf.constant(5.0)
```

```
b = tf.constant(6.0)
```

```
c = a * b
```

Numpy	TensorFlow
<code>a = np.zeros((2,2)); b = np.ones((2,2))</code>	<code>a = tf.zeros((2,2)), b = tf.ones((2,2))</code>
<code>np.sum(b, axis=1)</code>	<code>tf.reduce_sum(a, reduction_indices=[1])</code>
<code>a.shape</code>	<code>a.get_shape()</code>
<code>np.reshape(a, (1,4))</code>	<code>tf.reshape(a, (1,4))</code>
<code>b * 5 + 1</code>	<code>b * 5 + 1</code>
<code>np.dot(a,b)</code>	<code>tf.matmul(a, b)</code>
<code>a[0,0], a[:,0], a[0,:]</code>	<code>a[0,0], a[:,0], a[0,:]</code>

Also supports convolution: `tf.nn.conv2d`, `tf.nn.conv3D`

Tensor neural network functions

- Easy to define operations on layers of networks
 - `relu(features, name=None)`
 - `bias_add(value, bias, data_format=None, name=None)`
 - `sigmoid(x, name=None)`
 - `tanh(x, name=None)`
 - `conv2d(input, filter, strides, padding)`
 - `conv1d(value, filters, stride, padding)`
 - `conv3d(input, filter, strides, padding)`
 - `conv3d_transpose(value, filter, output_shape, strides)`
 - `sigmoid_cross_entropy_with_logits(logits, targets)`
 - `softmax(logits, dim=-1)`
 - `log_softmax(logits, dim=-1)`
 - `softmax_cross_entropy_with_logits(logits, labels, dim=-1)`
- Each function created *knows its gradient*
- **Automatic Differentiation** is just **chain rule**
- But... lets start simple...

Tensor function evaluation

```
import tensorflow as tf
```

```
a = tf.constant(5.0)
```

```
b = tf.constant(6.0)
```

```
c = a*b
```

```
with tf.Session() as sess:
```

```
    print(sess.run(c))
```

```
    print(c.eval())
```

output = 30

- Easy to define operations on tensors
 - constant
 - variables
 - placeholders
- Nothing evaluated until you define a session and tell it to evaluate it
- Session defines configuration of execution
 - like GPU versus CPU

Computation Graph with Code

```
import tensorflow as tf
X = tf.placeholder()
y = tf.placeholder()
```

$$J(\mathbf{W}) = \frac{1}{N} \sum_i^N (y^{(i)} - \underbrace{(\mathbf{W} \cdot \mathbf{x}^{(i)} + \mathbf{b})}_{y_{pred}})^2$$

1. **Setup** Variables and computations

```
W = tf.Variable("weights", (1,num_features),
               initializer=tf.random_normal_initializer())
b = tf.Variable("bias", (1,),
               initializer=tf.constant_initializer(0.0))

y_pred = tf.matmul(X,W) + b
loss = tf.reduce_sum((y-y_pred)**2)/n_samples
```

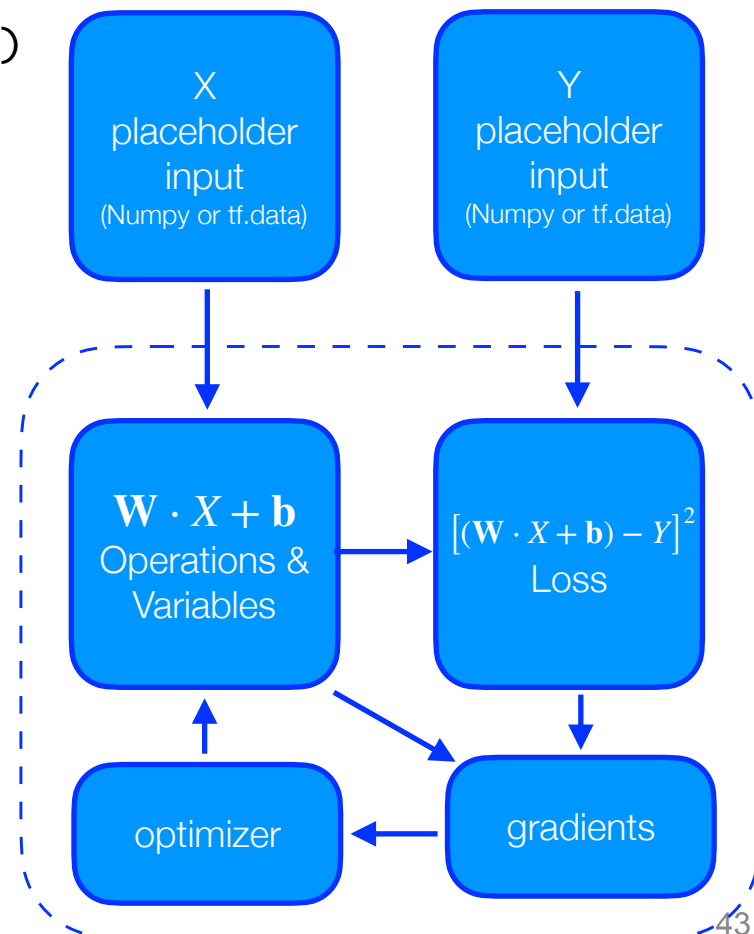
2. Add **optimization** operation to computation graph
Adjusts variables (W, b) to minimize loss with automatic differentiation

```
opt = tf.train.AdamOptimizer()
opt_operation = opt.minimize(loss)

with tf.Session() as sess:
    sess.run(tf.initialize_all_variables())
    sess.run([opt_operation],
             feed_dict={X:X_numpy, y: y_numpy})
```

3. **Run graph operation** once, → one optimization update on all variables

<http://www.datasciencecentral.com/profiles/blogs/google-open-source-tensorflow>



Tensorflow Mini-batching

```
opt = tf.train.AdamOptimizer()
opt_operation = opt.minimize(loss) Define reusable operation, → one optimization update

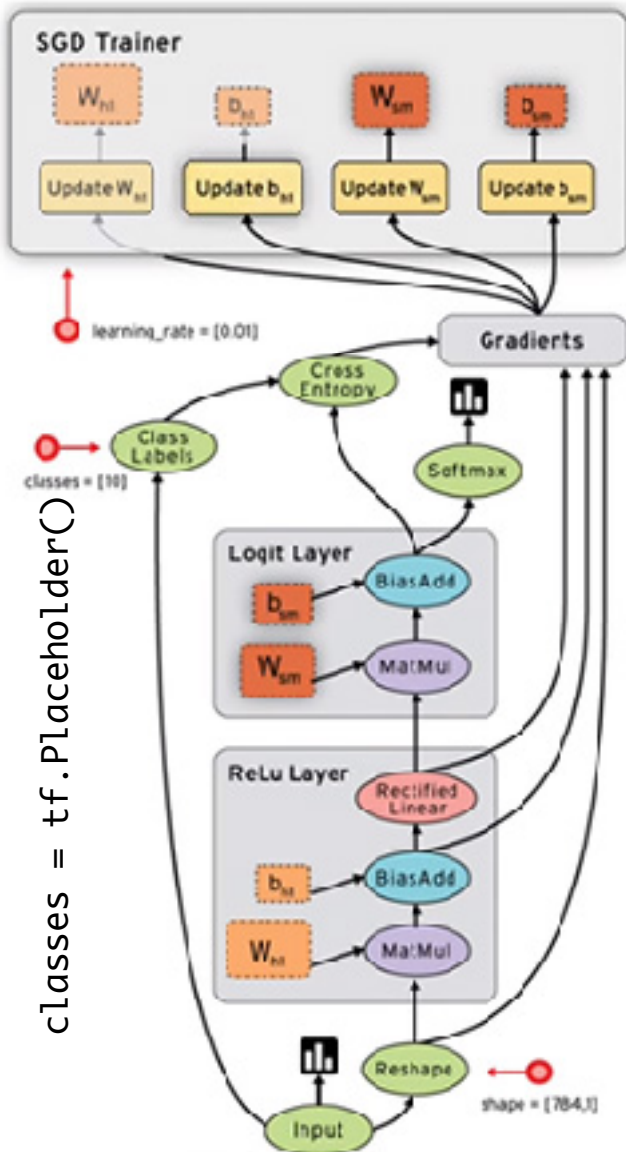
with tf.Session() as sess:
    sess.run(tf.initialize_all_variables())

    for _ in range(500):
        indices = np.random.choice(n_samples, batch_size) Get random
        X_batch, y_batch = X_numpy[indices], y_numpy[indices] training batch

        _, loss_val = sess.run([opt_operation, loss], Run optimization
                               feed_dict={X:X_batch, y:y_batch}) and return loss
```

- Example shown is **graph execution**
 - Build up computations and Execute computations when instructed
 - Makes it sometimes **hard to debug** but is **very fast**
- Alternative: **eager execution** (we won't cover this)

Computation Graph, Two Layer Network



```
Input = tf.placeholder() # size is 28x28
Input = tf.reshape(Input, [784,1])
classes = tf.placeholder()
```

```
W_sm = tf.Variable(...)
b_sm = tf.Variable(...)
W_hl = tf.Variable(...)
b_hl = tf.Variable(...)
```

```
A_hl = tf.relu( tf.matmul(Input,W_hl) + b_hl )
A_sm = tf.matmul(A_hl,W_sm) + b_sm
```

```
y_pr = tf.softmax(A_sm)
loss = tf.sparse_softmax_cross_entropy_with_logits(
    classes, A_sm )
```

```
opt = tf.train.SGDOptimizer(learning_rate=0.01)
opt_operation = opt.minimize(loss)
```


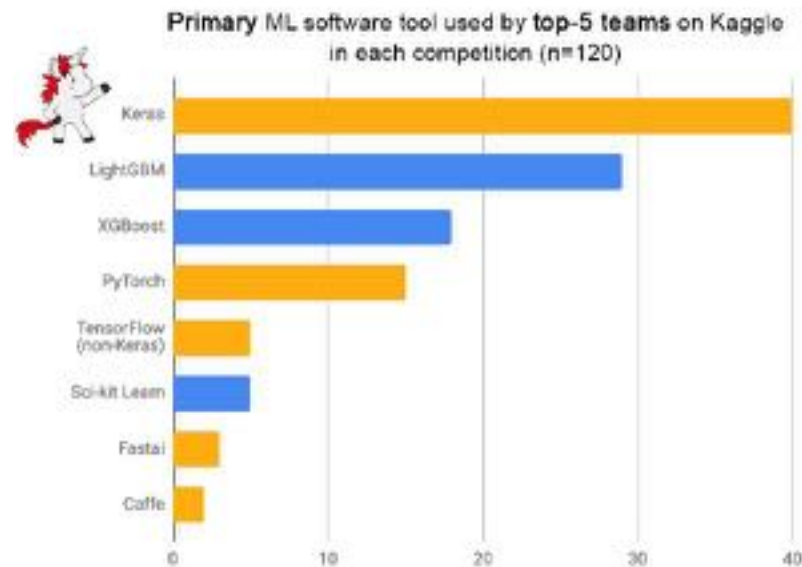

Tensorflow Simplification

- **Self Test:** Can the syntax be simplified?
 - (A) **Yes**, we could write a generic mini-batch optimization computation graph, then use it for arbitrary graph instructions
 - (B) **Yes**, but we need to learn the Keras API, which can be mixed with tensorflow operations
 - (C) **Yes**, but we need to understand how to access the gradients to apply them, a lot like PyTorch
 - (D) **All of the above**


Keras Programming Interfaces

- **Keras Sequential API**
 - great for simple, feed forward models
- **Keras Functional API**
 - build models through series of nested functions
 - each “function” represents an operation in the NN
- **Keras Classes (Inheritance)**
 - good for more advanced functionality

```
from tensorflow import keras
```



Reinventing the MLP
Wheel



10. Keras Wide and Deep.ipynb

10. Keras Wide and Deep as TFData.ipynb

Make me slow down if I go too fast!!