

Lecture Notes for **Machine Learning in Python**



Professor Eric Larson
An Ongoing History of Convolutional Networks

History of Convolutional Neural Networks



Machine Learning 101

Types of CNN, 1988-1998




Yann LeCun
Heads Facebook
AI Team

- **LeNet-1** (1988)
 - ~2600 params, not many layers
- **LeNet-5** (1998)
 - 7 layers, gets excellent MNIST performance
- Major contribution, general structure:
 - conv=>pool=>non-linearity=> ...=>MLP

avg

tanh or sigmoid



CNN History

- List of major breakthroughs from 1998 through 2010 in convolutional networks:



- 2010



Types of CNN, 2010



Dan Ciresan

AI Researcher
IDSA, Switzerland

- **Ciresan Net**
- Publishes code for running CNN via GPU
 - Subsequently wins 5 international competitions
 - from stop signs => cancer detection
- Major contribution: NVIDIA parallelized training algorithms


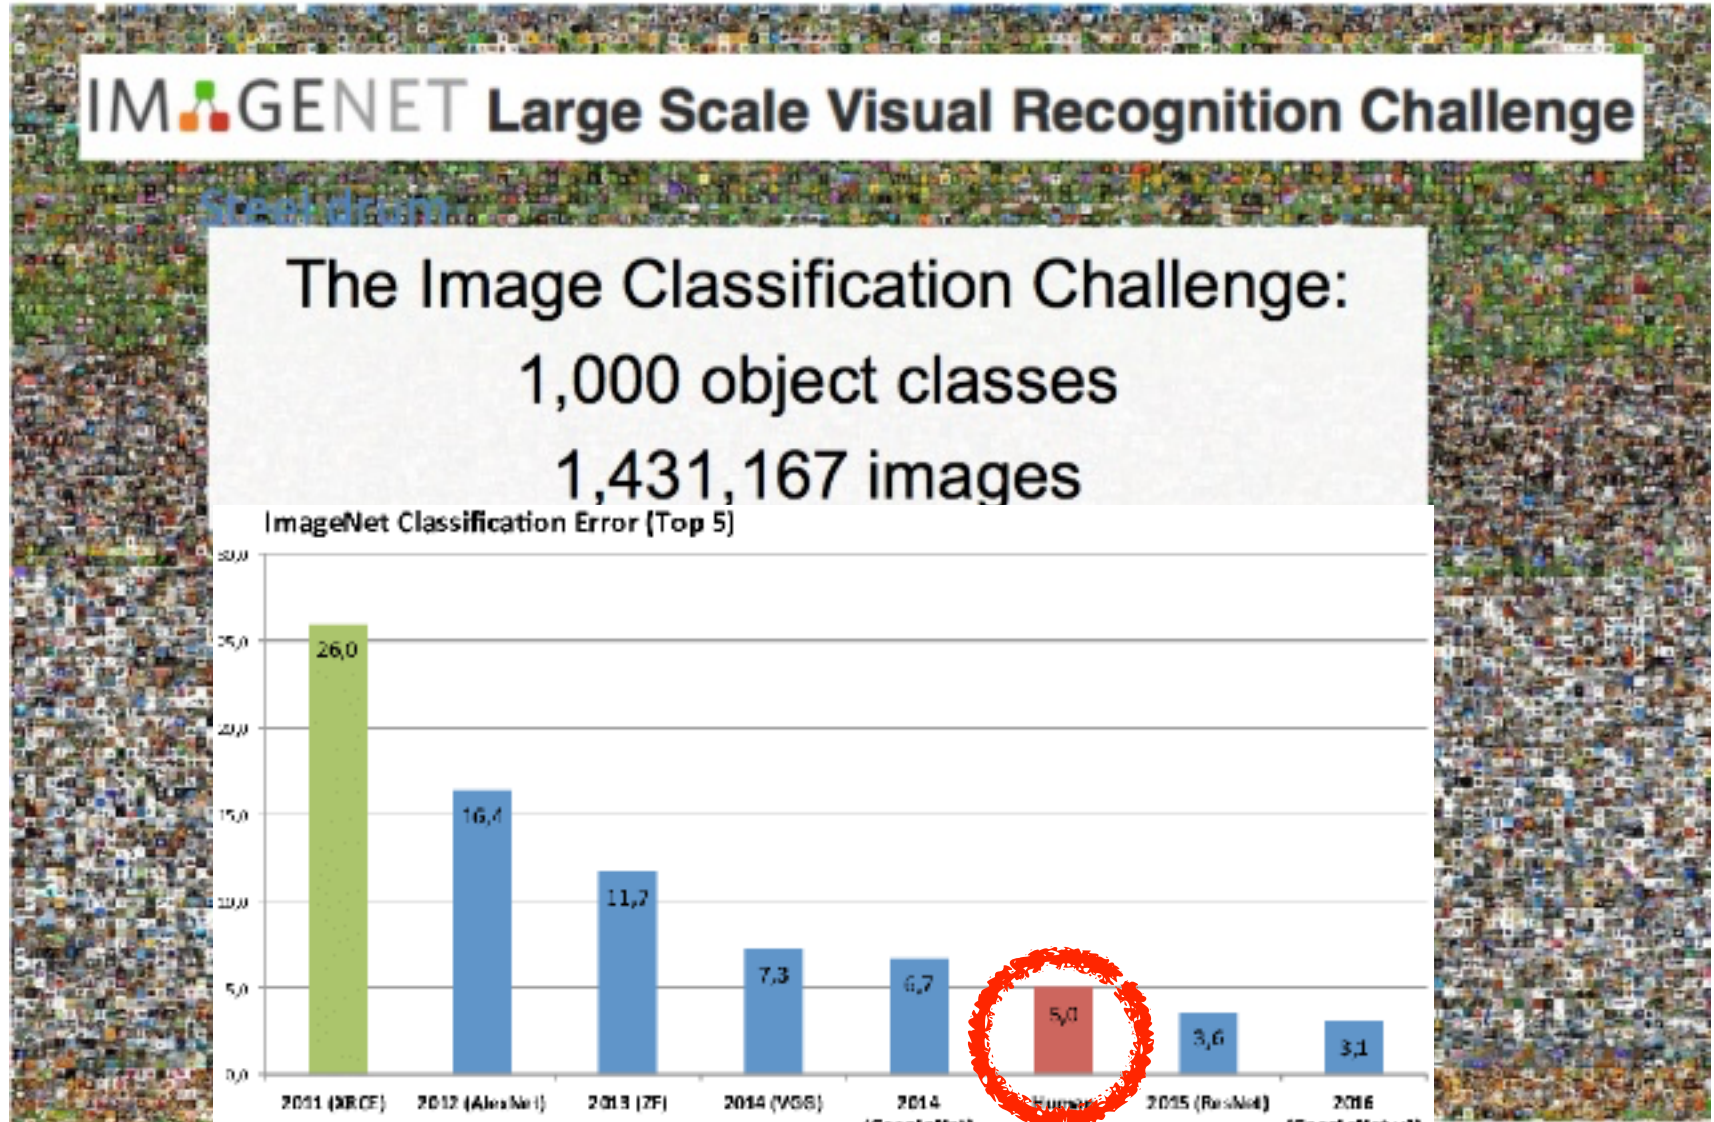


Figure 2: Forward propagation: a) mapping of kernel 1 grid onto the padded weight matrix; b) mapping the kernel 2 grid onto the partial dot products matrix; c) output of forward propagation.


ImageNet Competition (2010-2016)



https://www.researchgate.net/figure/Winner-results-of-the-ImageNet-large-scale-visual-recognition-challenge-LSVRC-of-the_fig7_324476862

<https://www.slideshare.net/nmhkahn/case-study-of-convolutional-neural-network-61556303>


Types of CNN, 2012



Alex Krizhevsky

Google

- **AlexNet**, Hinton is mentor
 - wins ImageNet competition
- Major contributions:
 - dropout for regularization
 - systematic use of ReLU
 - data expansion
 - ***overlapping max pool***



AlexNet

FC 1000

FC 4096 / ReLU

FC 4096 / ReLU

Max Pool 3x3s2

Conv 3x3s1, 256 / ReLU

Conv 3x3s1, 384 / ReLU

Conv 3x3s1, 384 / ReLU

Max Pool 3x3s2

Local Response Norm

Conv 5x5s1, 256 / ReLU

Max Pool 3x3s2

Local Response Norm


Conv 11x11s4, 96 / ReLU

Warning



WeKnowMemes

Types of CNN, 2013



Karen Simonyan
and
Andrew Zisserman



- Oxford **VGG Net** (Visual Geometry Group)
- Major contributions:
 - small cascaded kernels
 - way more layers (19 versus ~7)
 - “emulates” biology “better”
 - trained on NVIDIA GPUs for 2-3 weeks


ConvNet Configuration					
A	A-LRN	B	C	D	E
11 weight layers	11 weight layers	13 weight layers	16 weight layers	16 weight layers	19 weight layers
input (224 × 224 RGB image)					
conv3-64	conv3-64 LRN	conv3-64 conv3-64	conv3-64 conv3-64	conv3-64 conv3-64	conv3-64 conv3-64
maxpool					
conv3-128	conv3-128	conv3-128 conv3-128	conv3-128 conv3-128	conv3-128 conv3-128	conv3-128 conv3-128
maxpool					
conv3-256 conv3-256	conv3-256 conv3-256	conv3-256 conv3-256	conv3-256 conv3-256 conv1-256	conv3-256 conv3-256 conv3-256	conv3-256 conv3-256 conv3-256 conv3-256
maxpool					
conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512 conv1-512	conv3-512 conv3-512 conv3-512	conv3-512 conv3-512 conv3-512 conv3-512
maxpool					
conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512 conv1-512	conv3-512 conv3-512 conv3-512	conv3-512 conv3-512 conv3-512 conv3-512
maxpool					
FC-4096					
FC-4096					
FC-1000					
soft-max					

Table 2: Number of parameters (in millions).


Network	A,A-LRN	B	C	D	E
Number of parameters	133	133	134	138	144

- Network in Network **NiN**
 - or MLPConv


Min Lin^{1,2}, Qiang Chen², Shuicheng Yan²
¹Graduate School for Integrative Sciences and Engineering
²Department of Electronic & Computer Engineering
National University of Singapore, Singapore
{linmin, chenqiang, eleyans}@nus.edu.sg



(a) Linear convolution layer



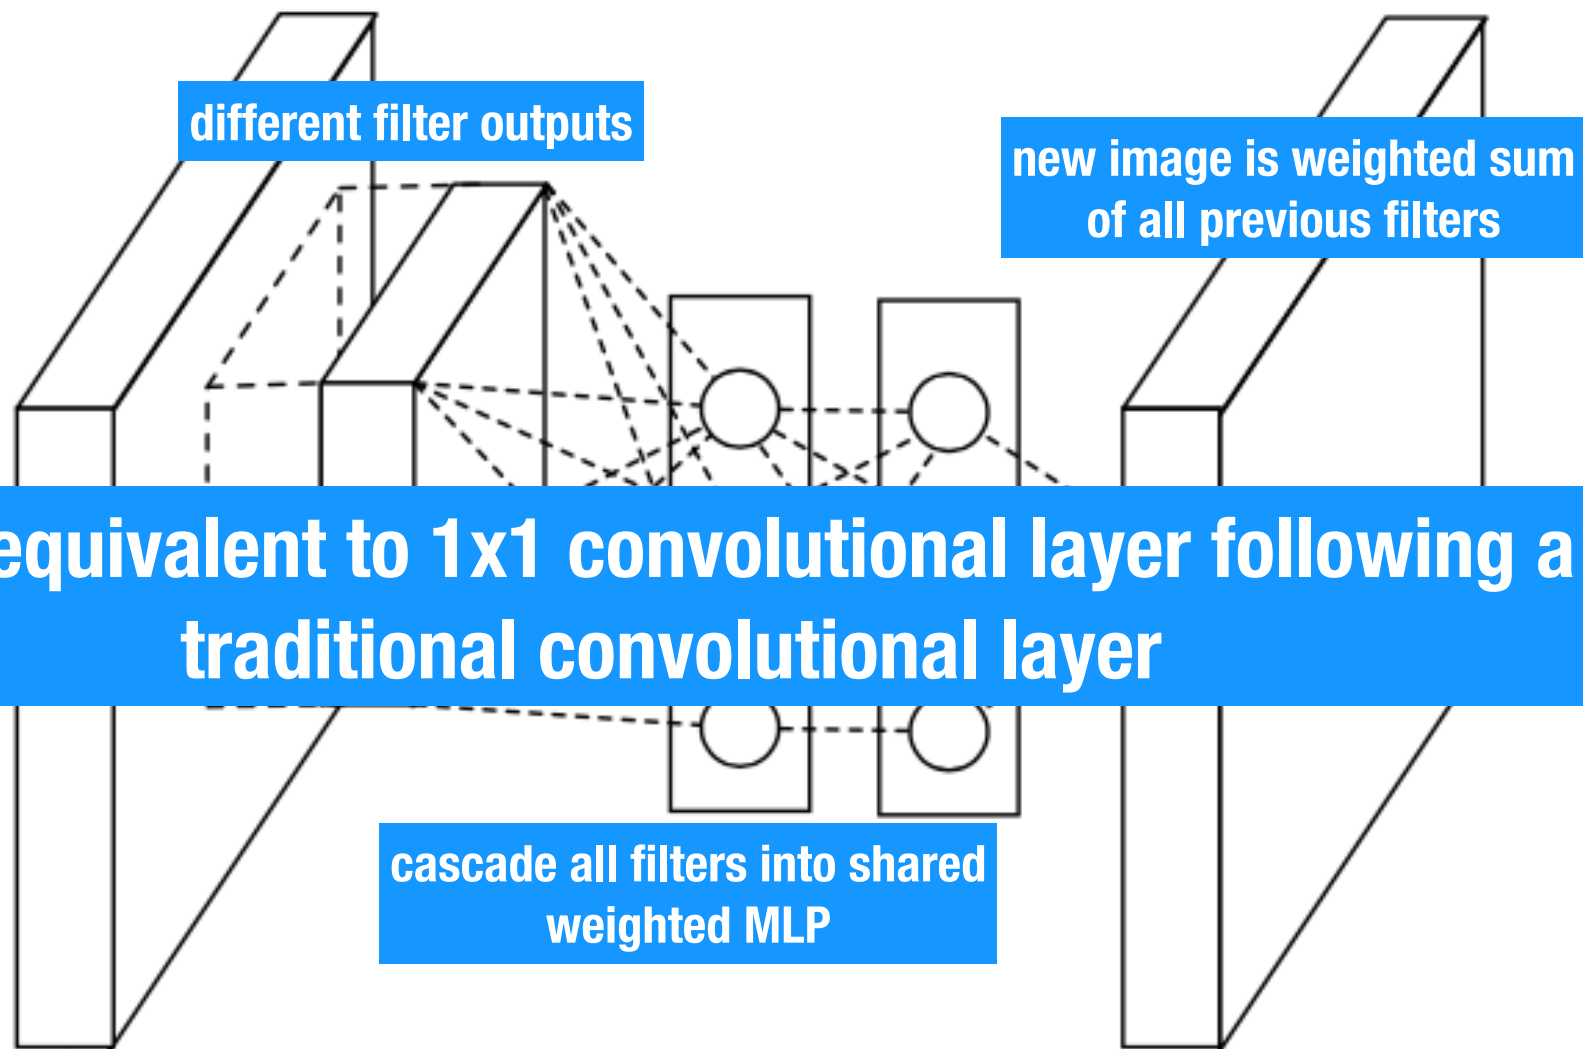
(b) Mlpconv layer



<https://arxiv.org/pdf/1312.4400.pdf>

Types of CNN, 2014


- Network in Network



NiN, expanded view

J and $M \gg K$ and L

Common Choice: $J=M$ and $K=L$




Convolve with K , 1×1 Filters

Equivalently: each new channel is weighted sum of convolutions complete control of channels size



one 1×1 filter input filter




one full filter, $k \times k$



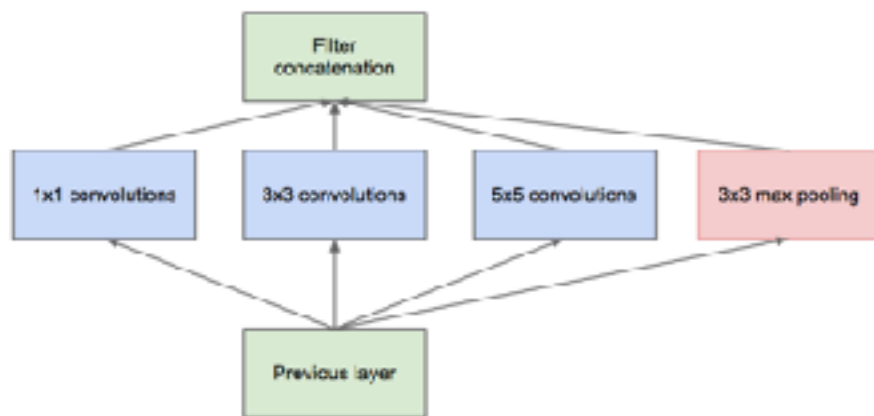
optional one 1×1 filters to control output size

Structure of Each Tensor: Channels x Rows x Columns


Types of CNN, 2014



- **GoogLeNet**
 - or **Inception V1**
- Major contribution:
 - bottleneck layering
 - parallel NiN



(a) Inception module, naïve version



(b) Inception module with dimension reductions

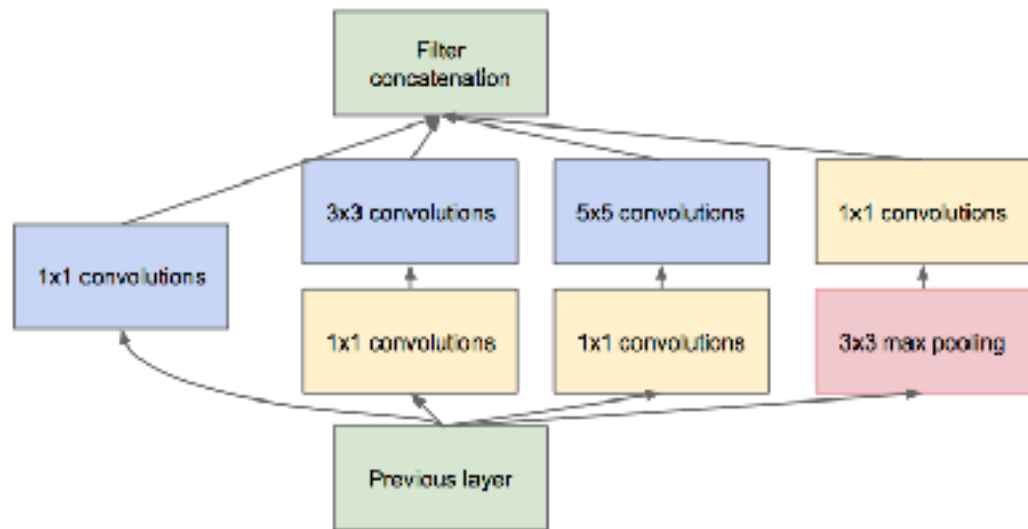



Figure 2: Inception module

Types of CNN, 2015 February and December






- **Inception V2**, Inception V1 with batch normalization
- **Inception V3:**
 - replace 5x5 with multiple 3x3




- Major Contributions:
 - “ensembles” not strictly sequential
 - “bio-plausible” with feedback

- ResNet**


- Parametric ReLU
- PReLU: adaptive trained slope



- NiN: triple bypass layer
 - similar to bottleneck




Residual Connection, expanded view




Back Propagation: Two paths, including one without ANY operations that cause the gradient to vanish...

How big are these networks?

How big are these networks?



34-layer residual



Transition Period in Convolutional Networks

- 2012 - 2017:
 - Add more layers! 🤪
 - How can we train it even deeper? 🧐
 - Can we run out of memory? Let's try! 🤔
- 2017-present:
 - How can we get similar performance with reduced parameters? 🤔
 - How should the number of parameters scale for competing resource? Is there an optimum scaling for a given set of resources? 📈

Types of CNN, 2017

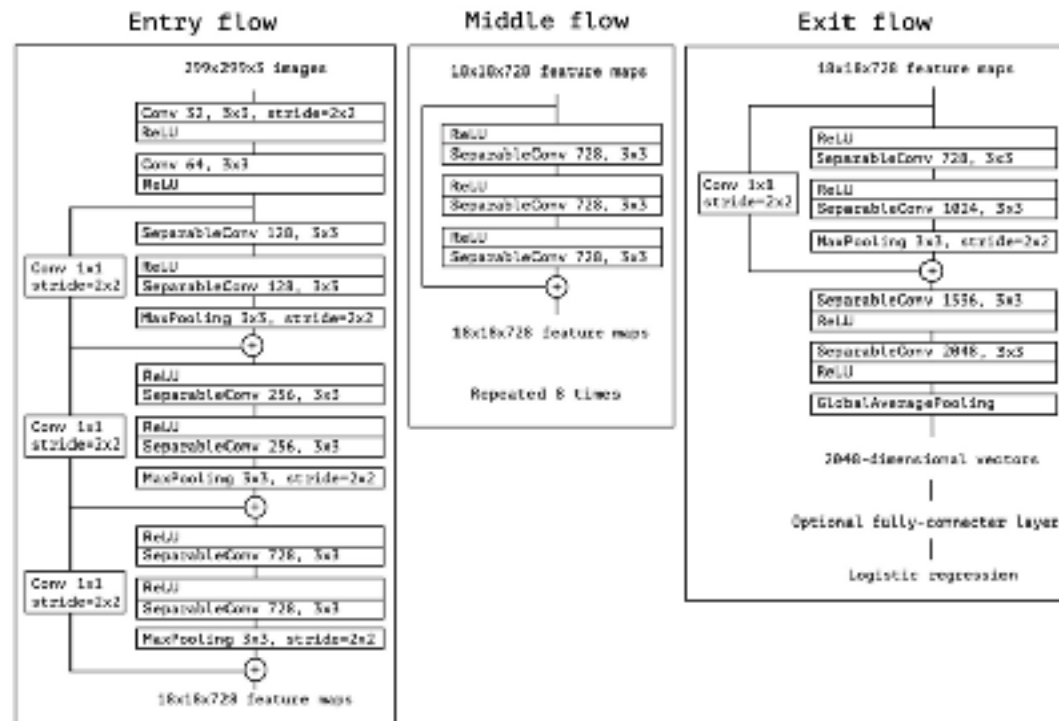
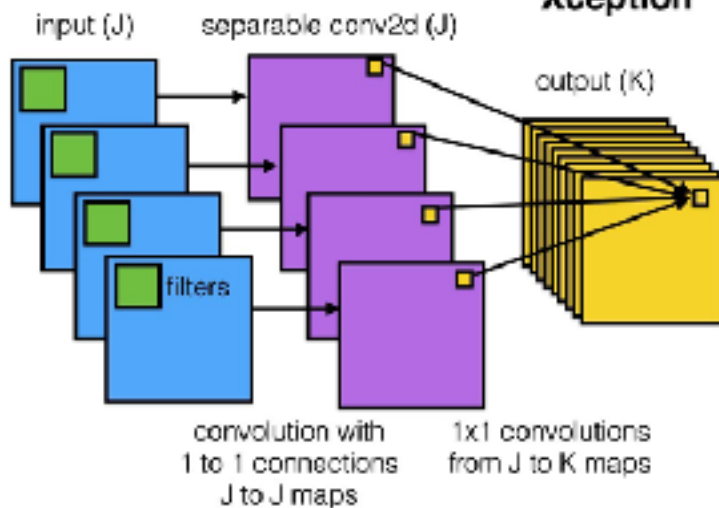
Xception · Major Contributions:

- combining branching / residual blocks
- separable convolutions (fewer trainable params)




Francois Chollet
Google

Xception




<https://arxiv.org/pdf/1610.02357.pdf> 63

Separable Convolution Explanation




Inputs, From Layer N-1
Num Channels = J



Filters, Layer N
Convolve Each Channel Separately

Trainable params:
Same as one filter in traditional convolution!




Concat Outputs
Num Channels = J

Perform K, 1x1 Traditional Convolutions
K Outputs


Trainable params: **K x J**

K Outputs



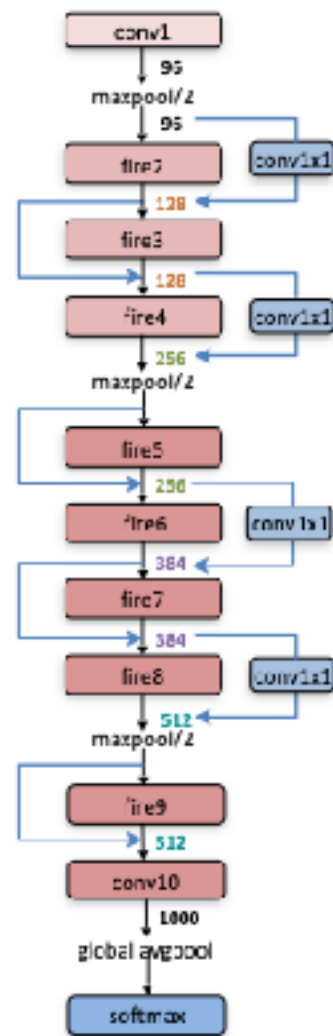
SqueezeNet (2018)

- Idea: squeeze and expand in each layer
 - Use mostly 1x1 filters
 - downsample later in network



$$SR = \frac{|F_{s1x1}|}{|F_{e1x1}| + |F_{e3x3}|}$$

$$PCT_{3x3} = \frac{|F_{e3x3}|}{|F_{e1x1}| + |F_{e3x3}|}$$



SQUEEZENET: ALEXNET-LEVEL ACCURACY WITH 50X FEWER PARAMETERS AND <0.5MB MODEL SIZE

Forrest N. Iandola¹, Song Han², Matthew W. Moskewicz¹, Khalid Ashraf¹, William J. Dally², Kurt Keutzer¹

¹DeepScale^{*} & UC Berkeley ²Stanford University

{forresti, moskewicz, kashraf, keutzer}@eecs.berkeley.edu

{songhan, dally}@stanford.edu

In paper:

- Good SR = 12.5% up to 100%
- Good PCT_{3x3} from 25% up to 100%

Efficient Net (2019)

Start with so

Observation 1 – Scaling up any width, depth, or resolution improve accuracy gain diminishes for bigger models

Observation 2 – In order to pursue efficiency, it is critical to balance all width, depth, and resolution during

Depth Scaling

Resolution Scaling: If we use larger resolution

depth: $d = \alpha^\phi$

width: $w = \beta^\phi$

res.: $r = \gamma^\phi$

s.t. $\alpha \cdot \beta^2 \cdot \gamma^2 \approx 2$

$\alpha, \beta, \gamma \geq 1$

ϕ user specified scaling coefficient

$\alpha = 1.2$

$\beta = 1.1$

$\gamma = 1.15$




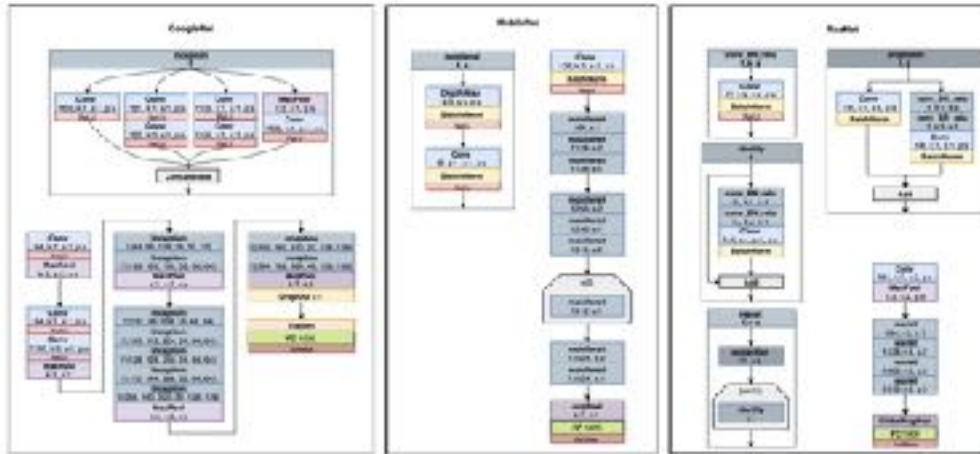
Figure 8. Scaling Up EfficientNet-B0 with Different Methods.

where α, β, γ are constants that can be determined by a small grid search. Intuitively, ϕ is a user-specified coefficient that controls how many more resources are available for model scaling, while α, β, γ specify how to assign these extra resources to network width, depth, and resolution re-

optimal values found in paper!

Even more Convolutional
Neural Networks
...in TensorFlow
...with Keras

Self Guided Demo



12. More Advanced CNN Techniques as TFData.ipynb

Next Time:

- Intro to Sequential Neural Network Architectures
 - Word Embeddings, 1D CNNs, Transformers
 - Ethics by Case Study